



Containers & Virtualization

Emmanuel Druon



Global organization

- Module #1
 - Virtualization: VM / Clusters / Docker / Network / Storage
 - Cloud concepts
 - Presentation of different Cloud offers
- Module #2: Instances & Networking
 - Lab: Create a simple EC2 instance
 - Lab: Build a network infrastructure with a VPC
- Module #3: Cloud Storage & Containers
 - Introduction to the main types of storage in AWS
 - Lab: S3 Lab
 - Lab: EBS Lab
 - A short introduction to Docker containers
 - Lab: Create a Web Service based on an “nginx” container
- Module #4: A short introduction to DevOps
 - Lab: Working with Ansible



Group work

- In groups, try to define:
 - What cloud computing is?
 - Pros & Cons?
 - Types?
- Timing: 15 minutes



What is Cloud Computing?

“ Cloud computing is massively based on virtualization technology ”

BUT

Cloud computing is more than virtualization



Cloud Computing?

- Technically, Cloud Computing is based on proposing resource pools through:
 - Abstraction: virtualization,
 - Machine virtualization,
 - Network virtualization,
 - Storage virtualization.
 - Automation: orchestration,
 - Provision,
 - Deprovision,
 - Resize.



- “Classical” virtualization vs Cloud computing

- “Classical” virtualization:
 - Abstraction,
 - Administrator oriented,
 - Not self-service,
 - Poor elasticity.
- Cloud computing:
 - Abstraction,
 - Orchestration of resource pools,
 - Self-service.



Cloud computing definition

- National Institute of Standards and Technology's (NIST) definition:
 - *"Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction."*



Benefits of Cloud Computing

- Agility:
 - Possibility to quickly provision & deprovision resources,
- Resilience:
 - Load-balancing,
 - Duplication.
- Cost:
 - Pay for use,
- Security.



Virtualization in a nutshell



Virtualization in a nutshell

Machine virtualization



Machine Virtualization

- Two distinct models:
 - Full scale Virtual Machines,
 - Containers.

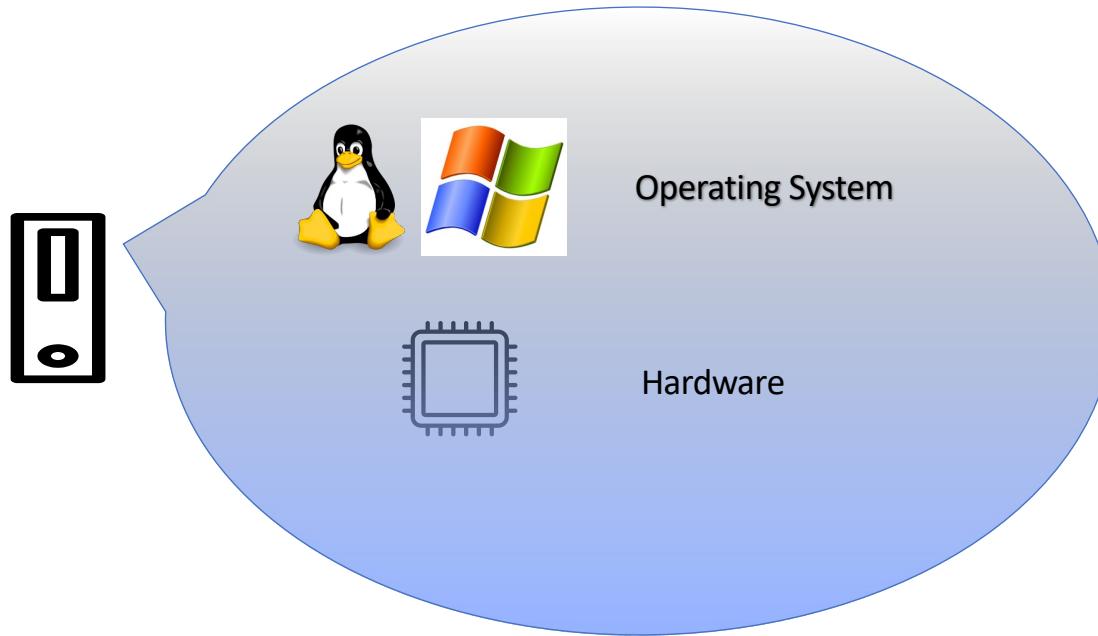


Full scale VMs

- Principles
 - The hardware is handled by a hypervisor,
 - Virtual hardware is software defined,
 - VMs are “executed” under the supervision of the hypervisor,
 - Each VM has its own OS,
 - VMs are isolated from one another.

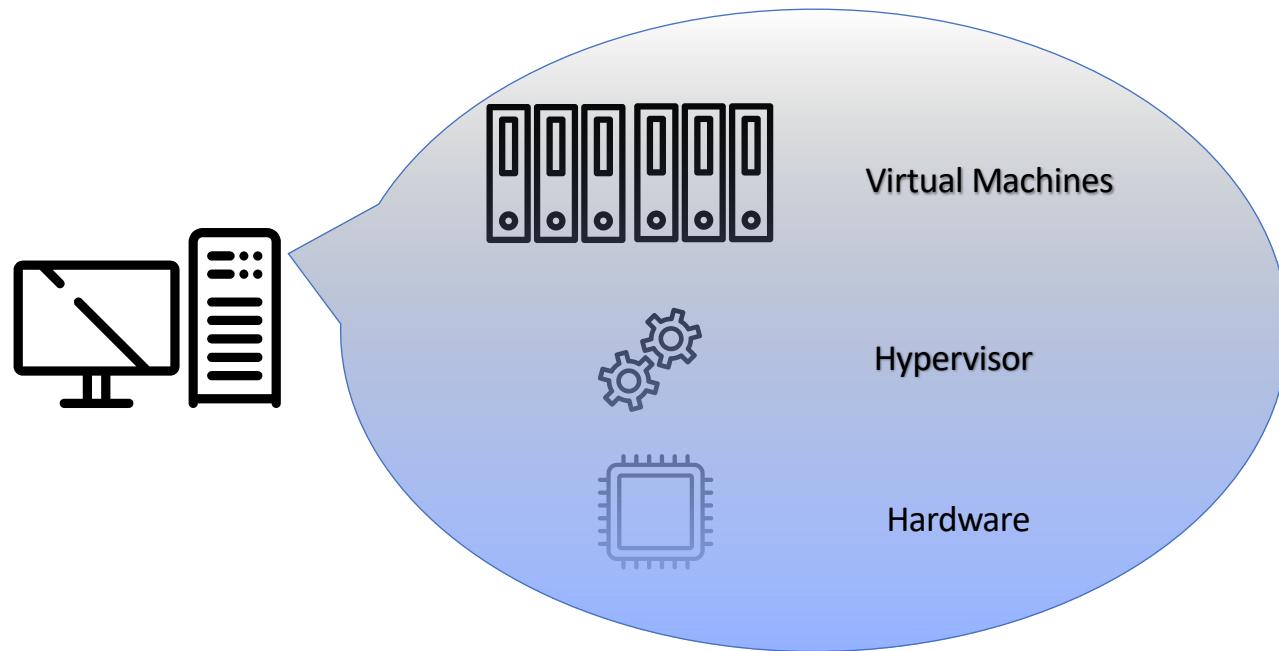


Full scale VMs



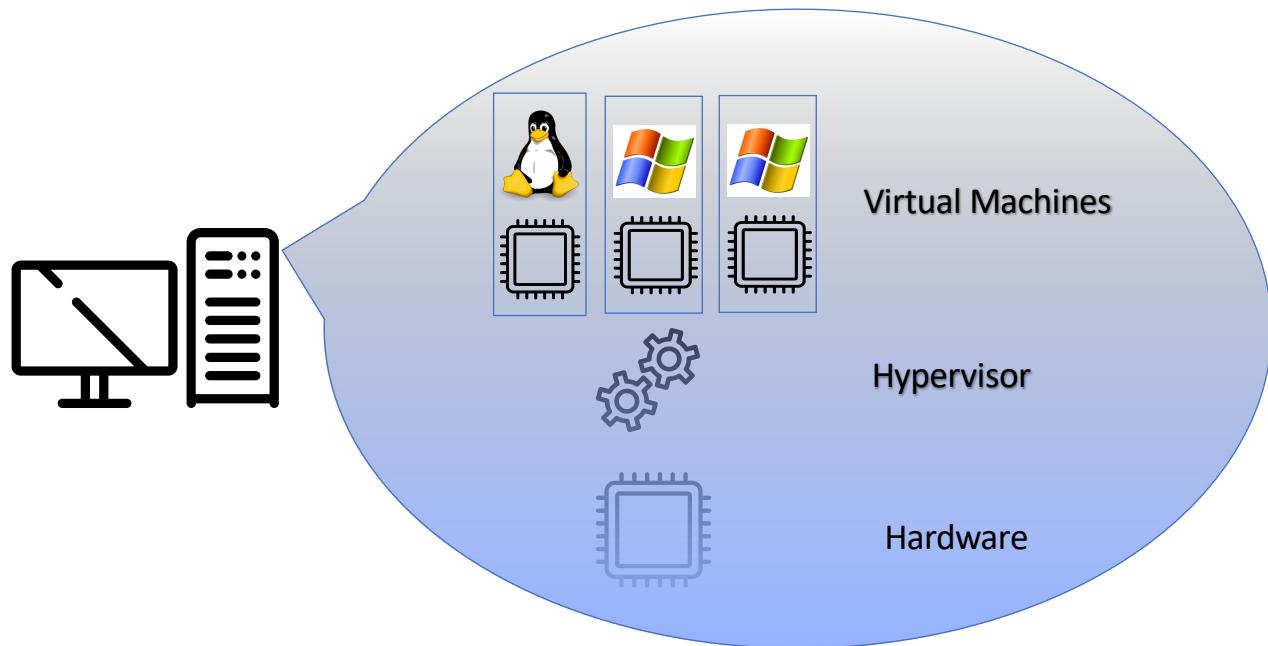


Full scale VMs





Full scale VMs





A VM configuration file example

```
[edruon@21-3910-001 temp % cat Debian\ 11.x\ console.vmx
usb.vbluetooth.startConnected = "TRUE"
firmware = "efi"
guestOS = "arm-debian11-64"
sound.autoDetect = "TRUE"
sound.virtualDev = "hdaudio"
cpuid.coresPerSocket = "1"
memsize = "2048"
nvme0.present = "TRUE"
nvme0:0.fileName = "Virtual Disk.vmdk"
nvme0:0.present = "TRUE"
sata0:1.deviceType = "cdrom-image"
sata0:1.fileName = "/Users/edruon/Downloads/debian-11.4.0-arm64-netinst.iso"
ethernet0.connectionType = "nat"
ethernet0.virtualDev = "e1000e"
ethernet0.generatedAddress = "00:0c:49:ef:5a:84"
```

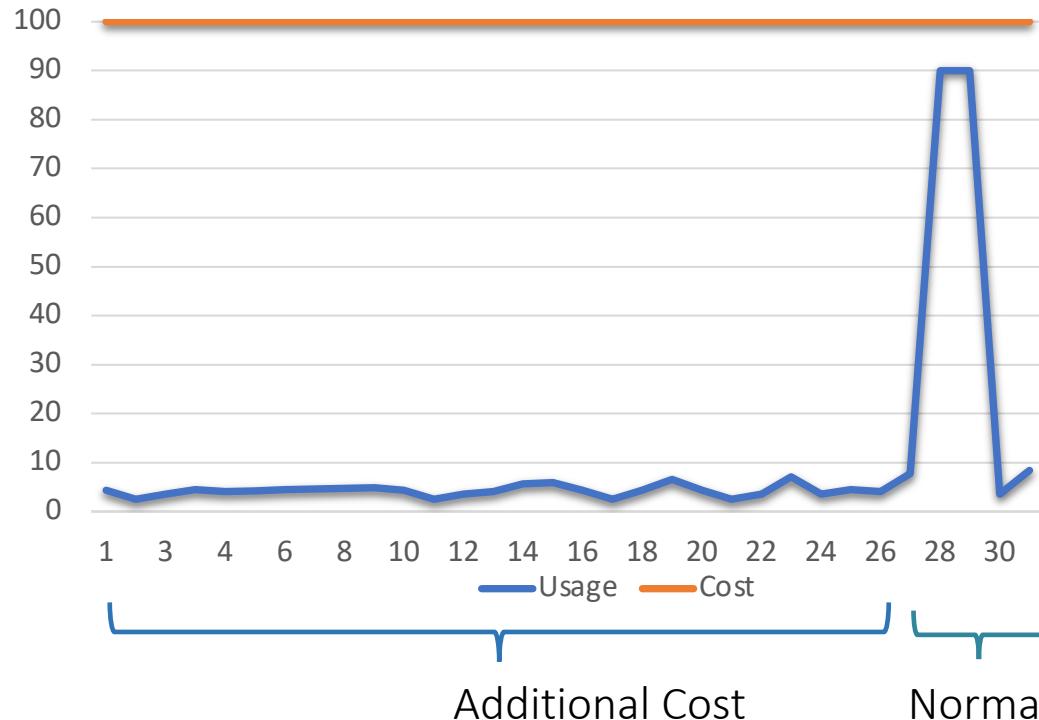


Great! I got it!

- But... Do we need virtualization?

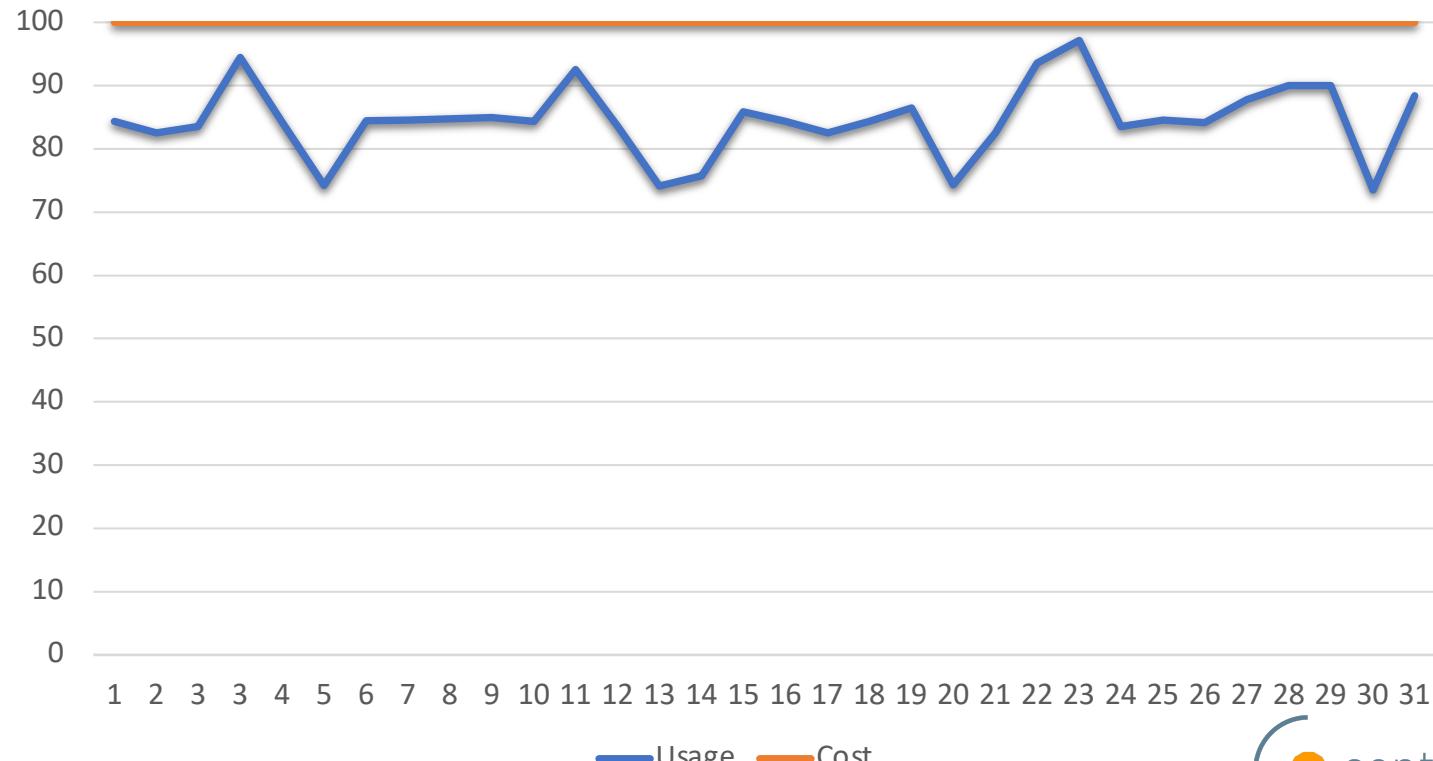


Why virtualization?





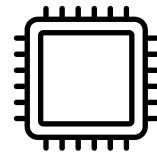
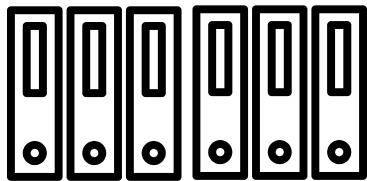
Why virtualization?



— Usage — Cost



Some disadvantages of virtualization



- Raw performance loss
- An extra software layer
- A possible unique point of failure
- Additional technical expertise required



Seriously?

- It's not that good in fact...



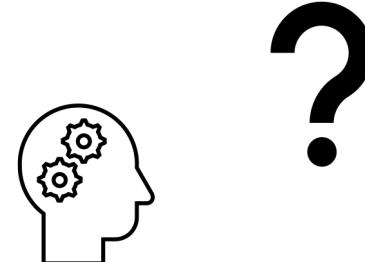
Seriously?

- It's not that good in fact... Or is it?



Before moving on... A simple question

What are the “real life components”
of a Virtual Machine?

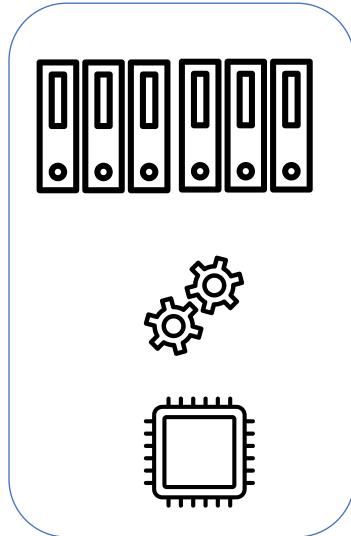




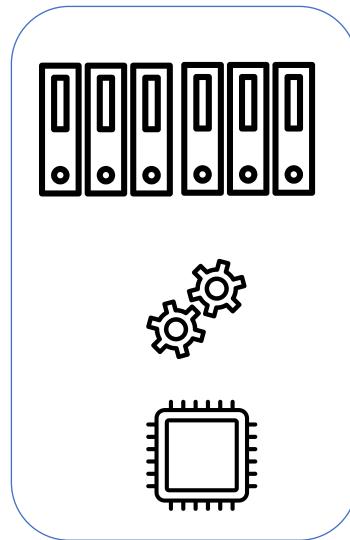
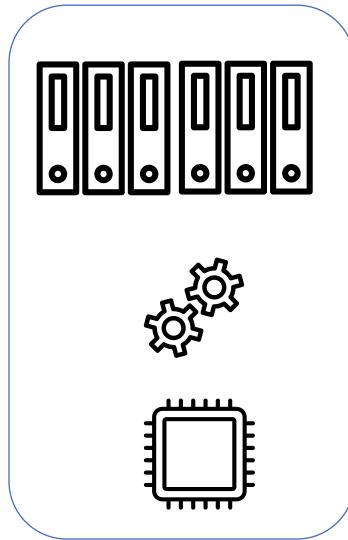
Before moving on... A simple question

- In real life, a VM is basically?
 - A set of files
 - A memory area on the server
 - Files represent the non-volatile part of a VM:
 - *.nvram* file saves the BIOS of the VM
 - *.vmx* and *.vmxf* keep the configuration
 - *.vmsd* files are used for snapshots
 - *.vmdk* files are the virtual disks used by the VM
 - Additionally, *.vmem* and *.vmss* files can be found when the VM is suspended or if it crashed. They save the memory pages used by the VM and the VM state when suspended
- Debian 11.x console.nvram
Debian 11.x console.vmsd
Debian 11.x console.vmx
Debian 11.x console.vmxr
Virtual Disk-s001.vmdk
Virtual Disk-s002.vmdk
Virtual Disk-s003.vmdk
Virtual Disk-s004.vmdk
Virtual Disk-s005.vmdk
Virtual Disk-s006.vmdk
Virtual Disk.vmdk
vmware-0.log

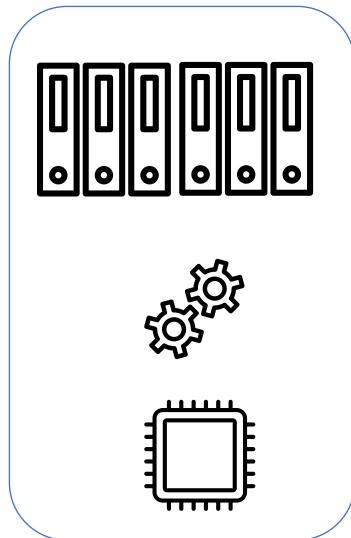
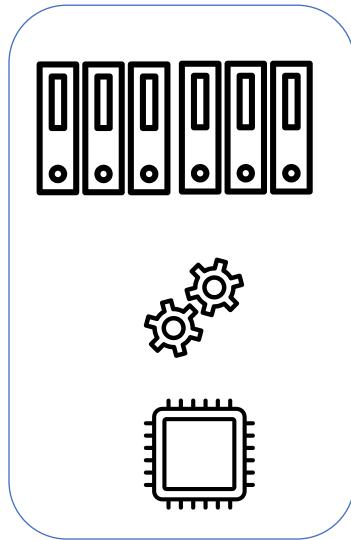
The cluster



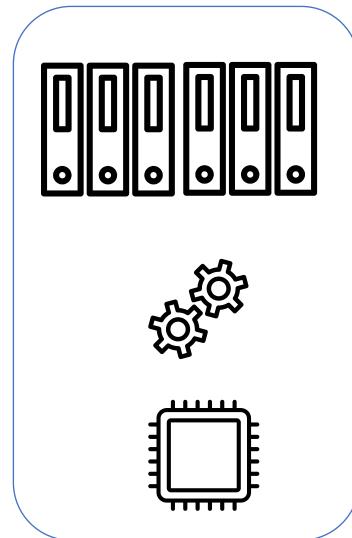
The cluster



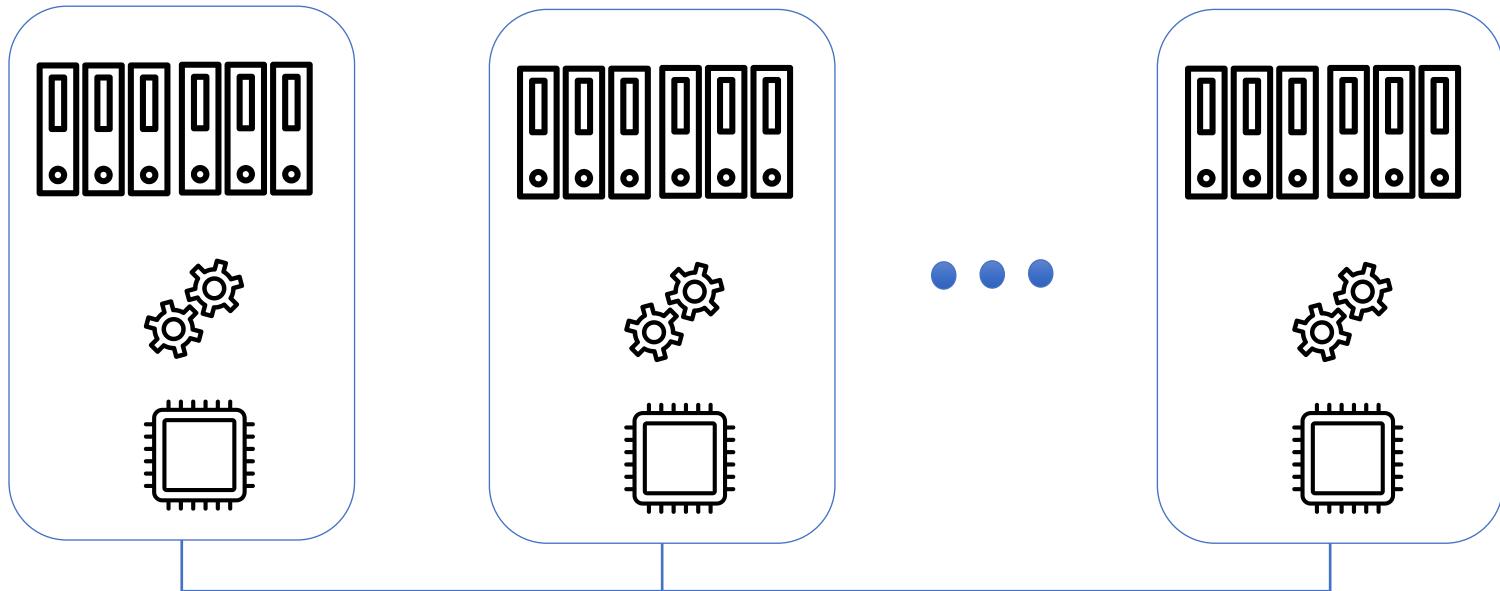
The cluster



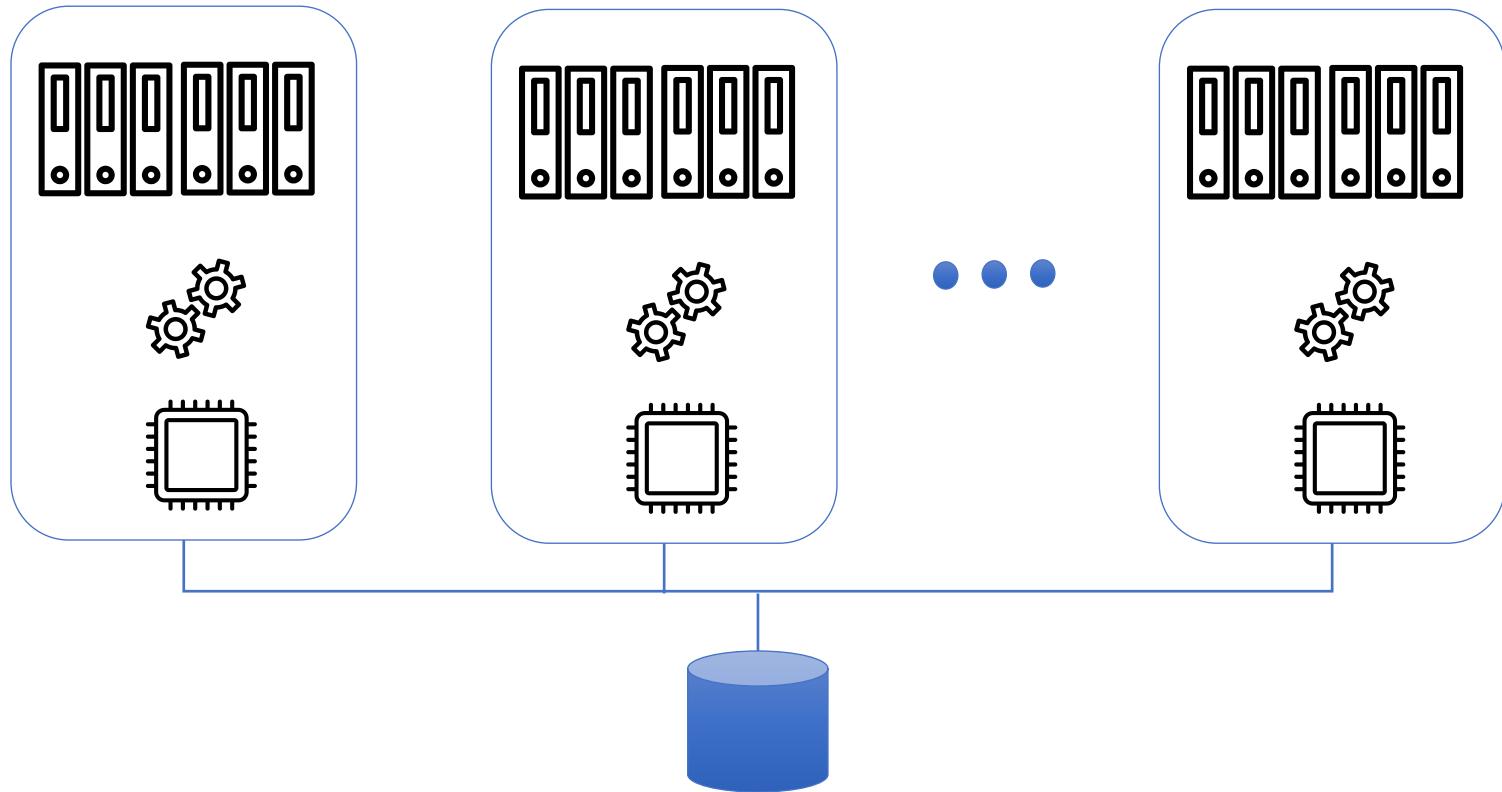
...



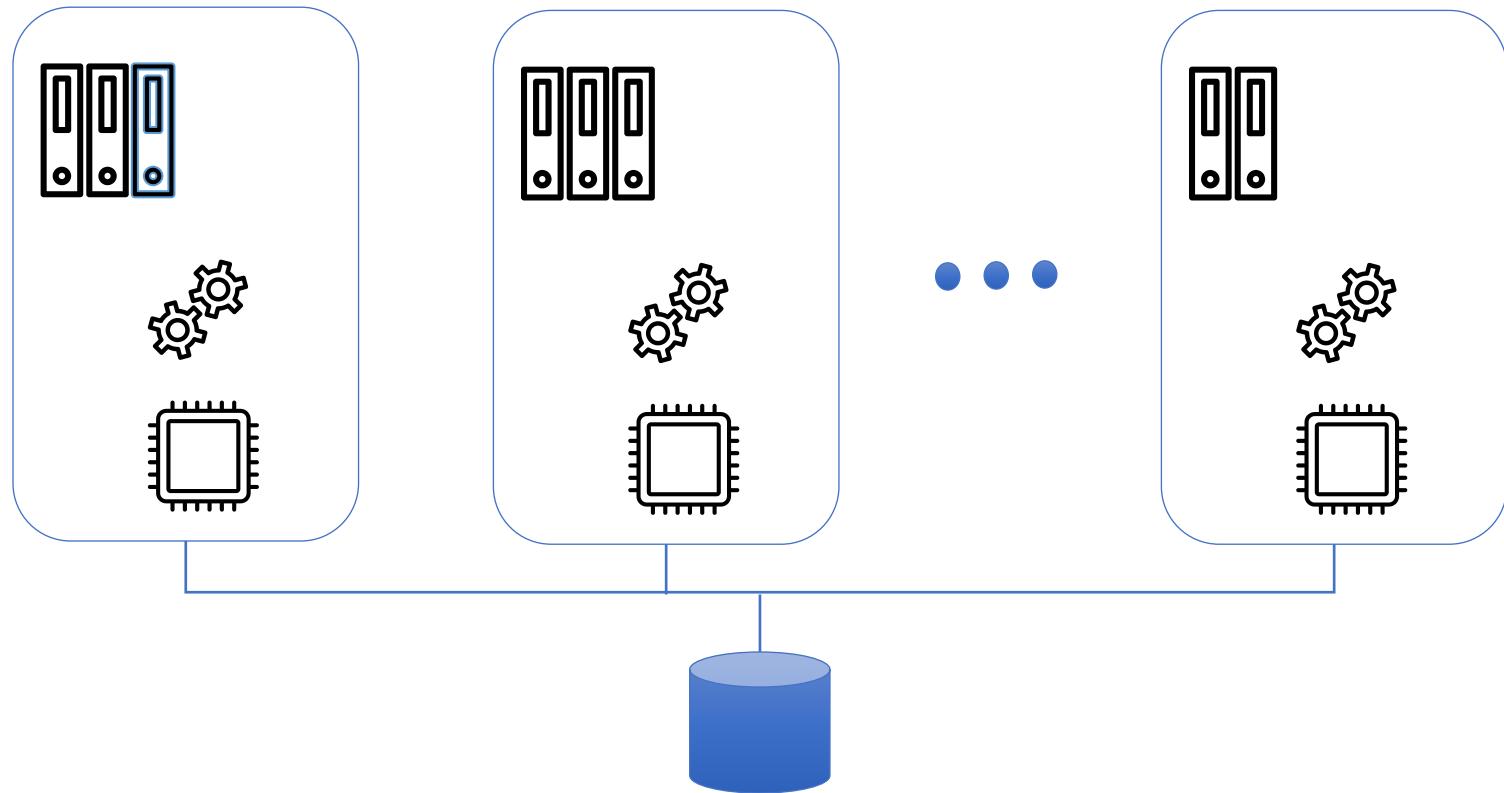
The cluster



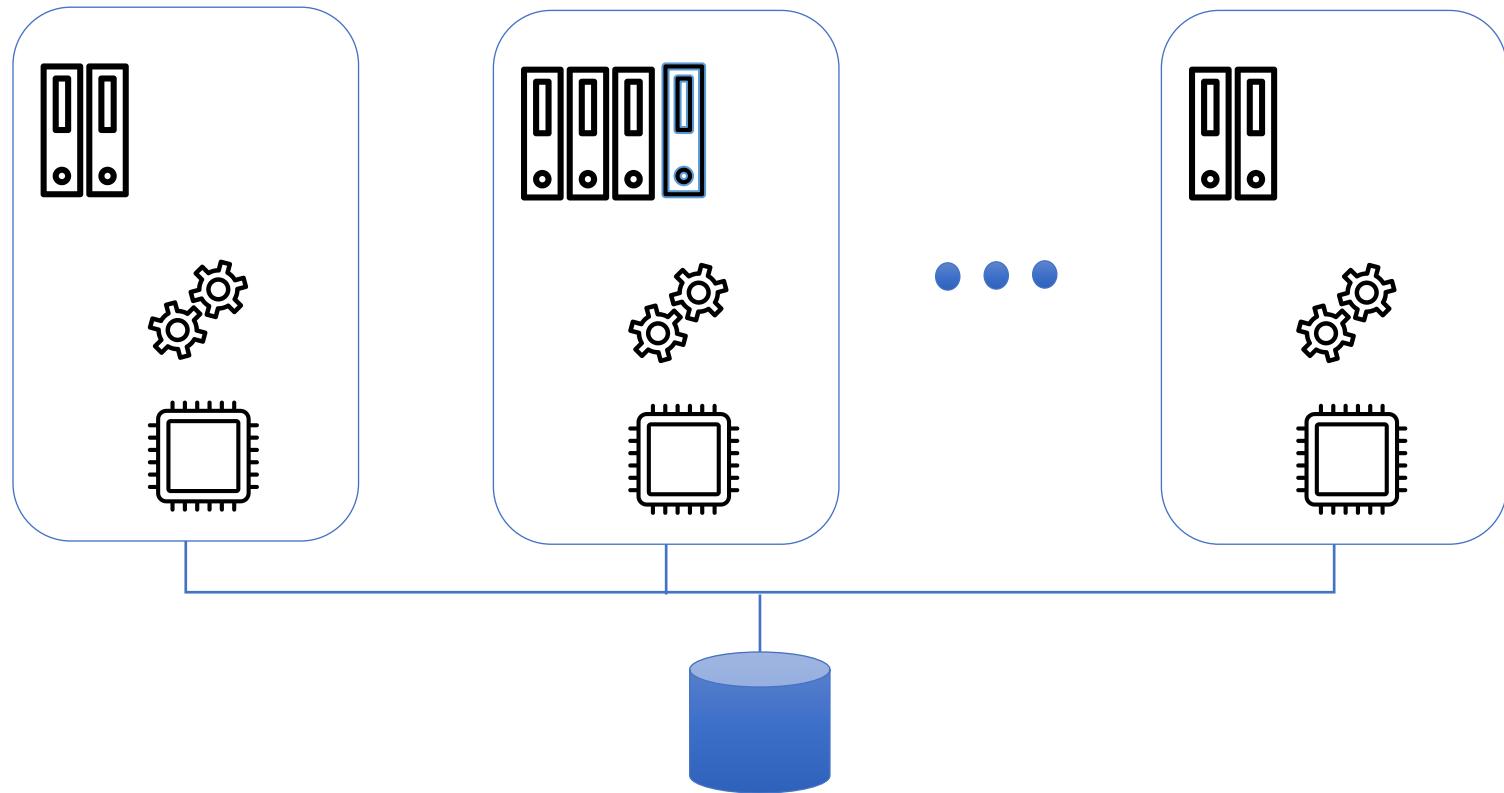
The cluster



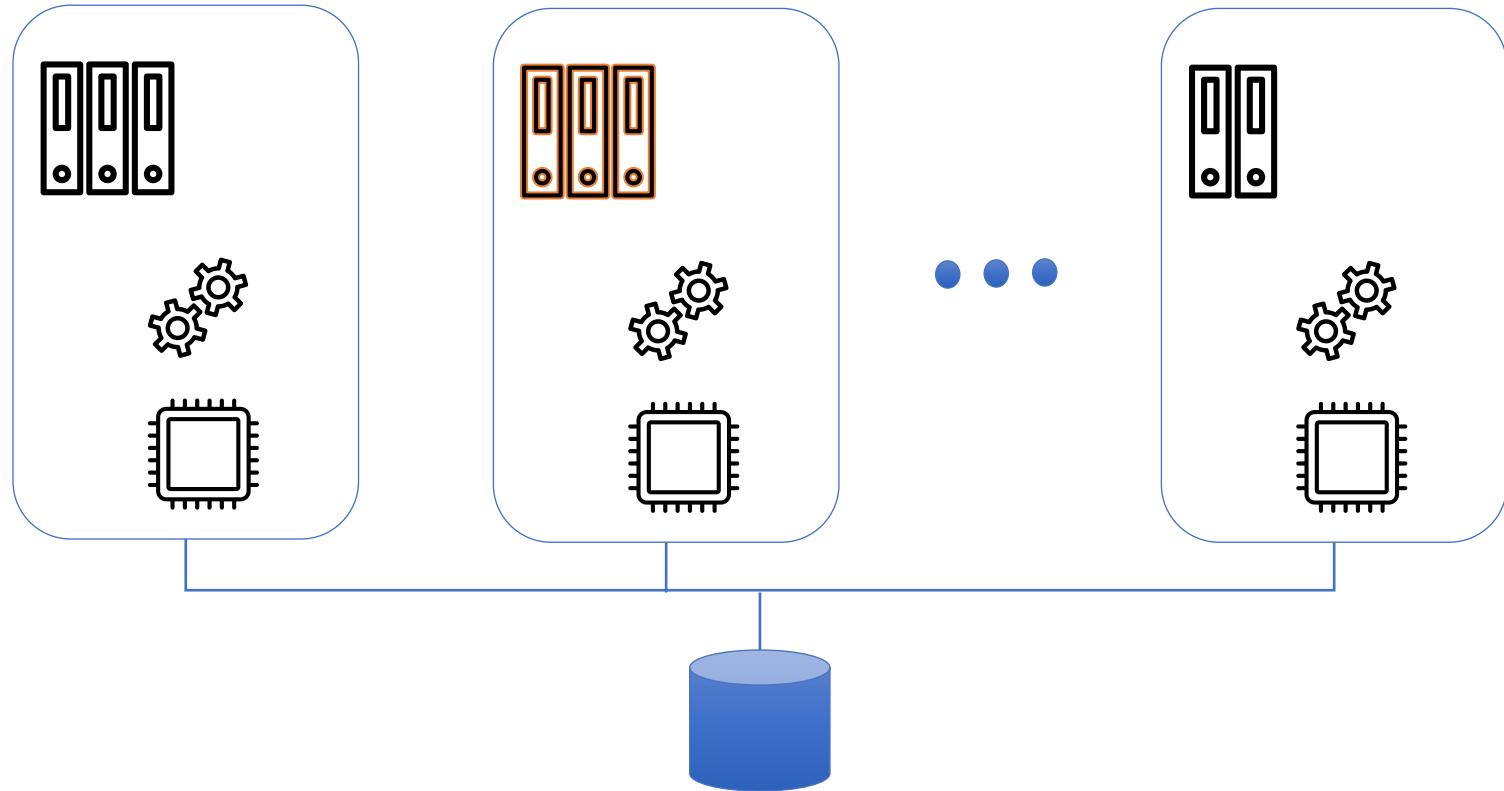
VM migration



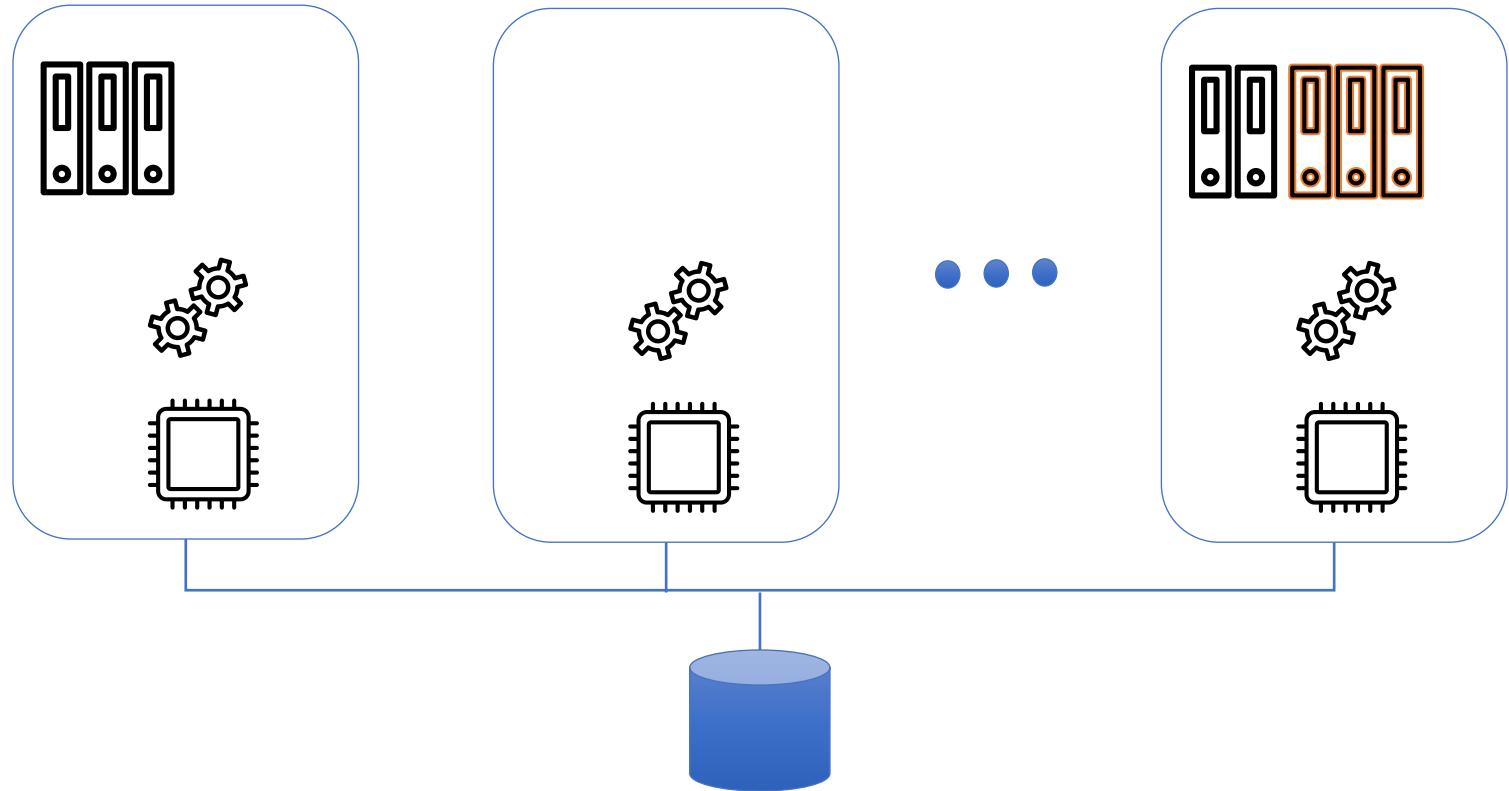
VM migration



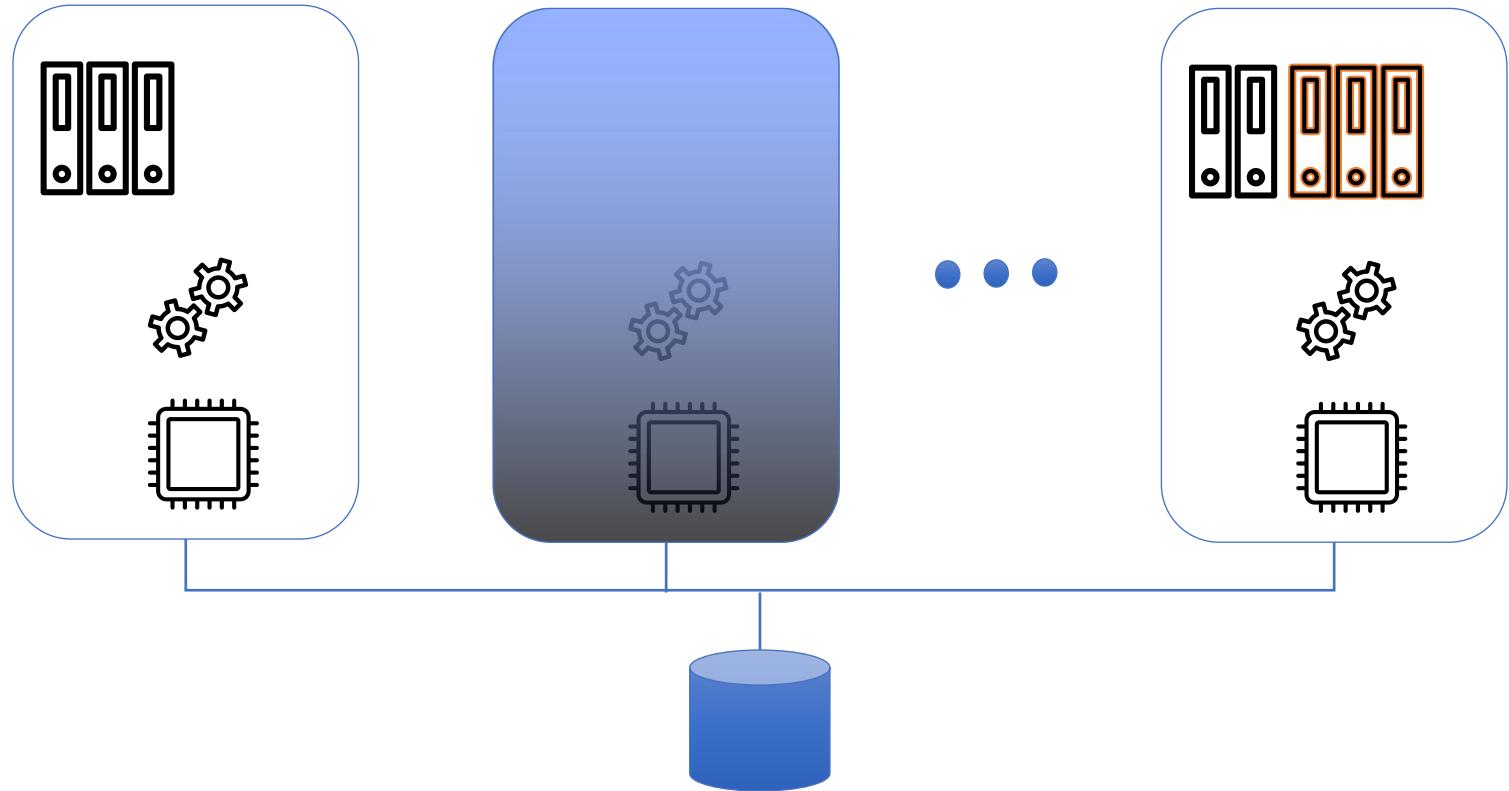
Physical server maintenance



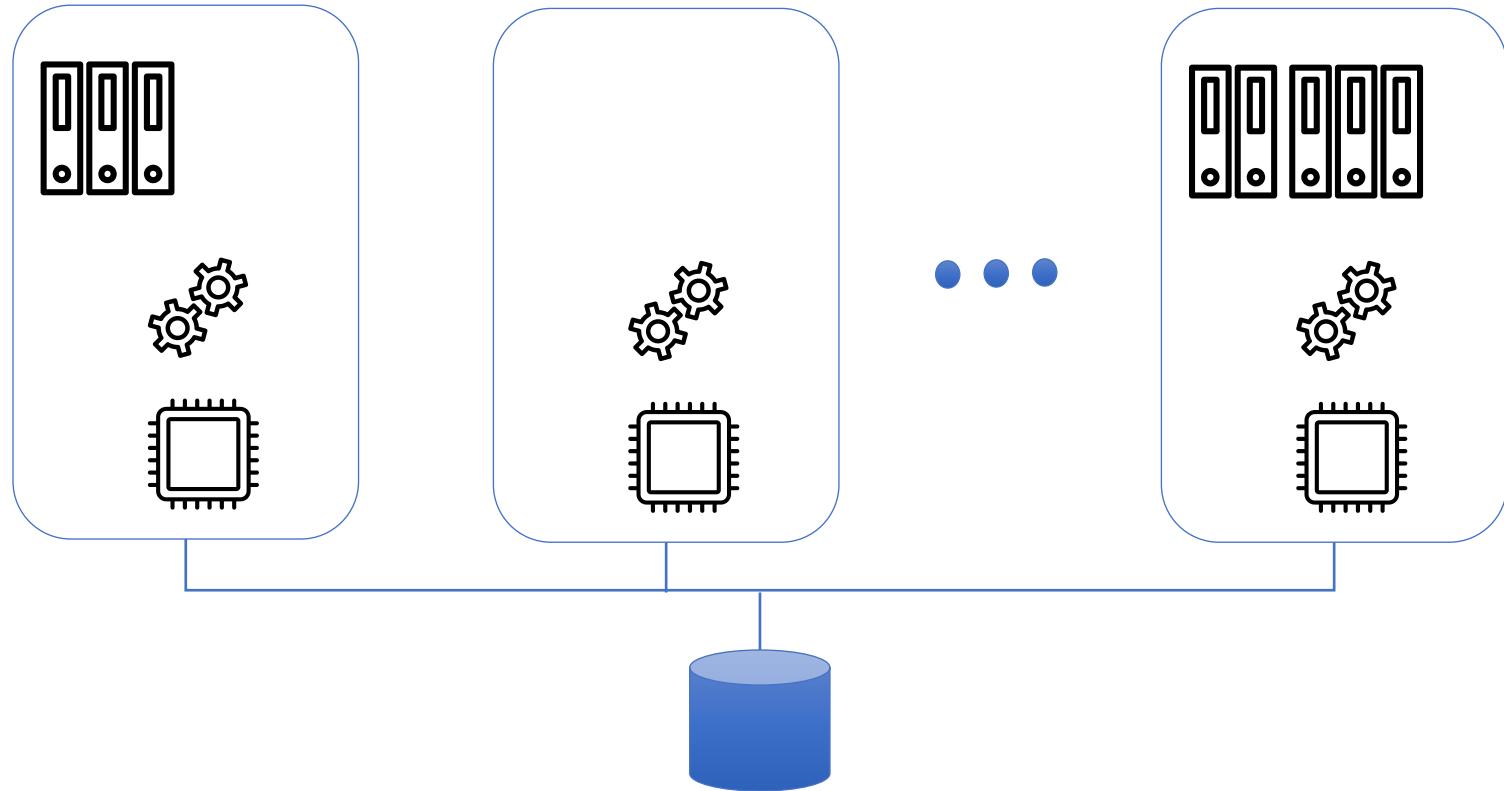
Physical server maintenance



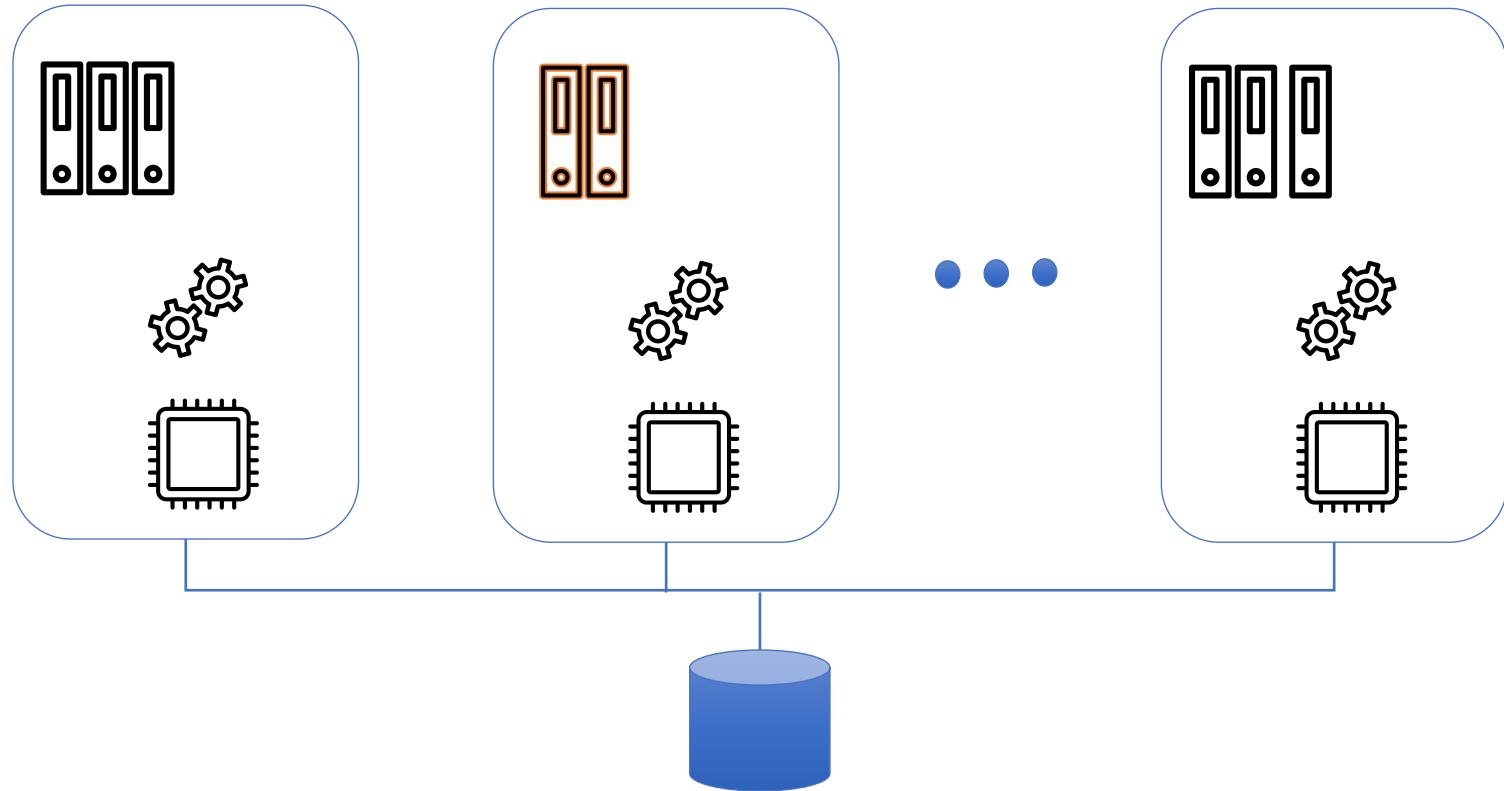
Physical server maintenance



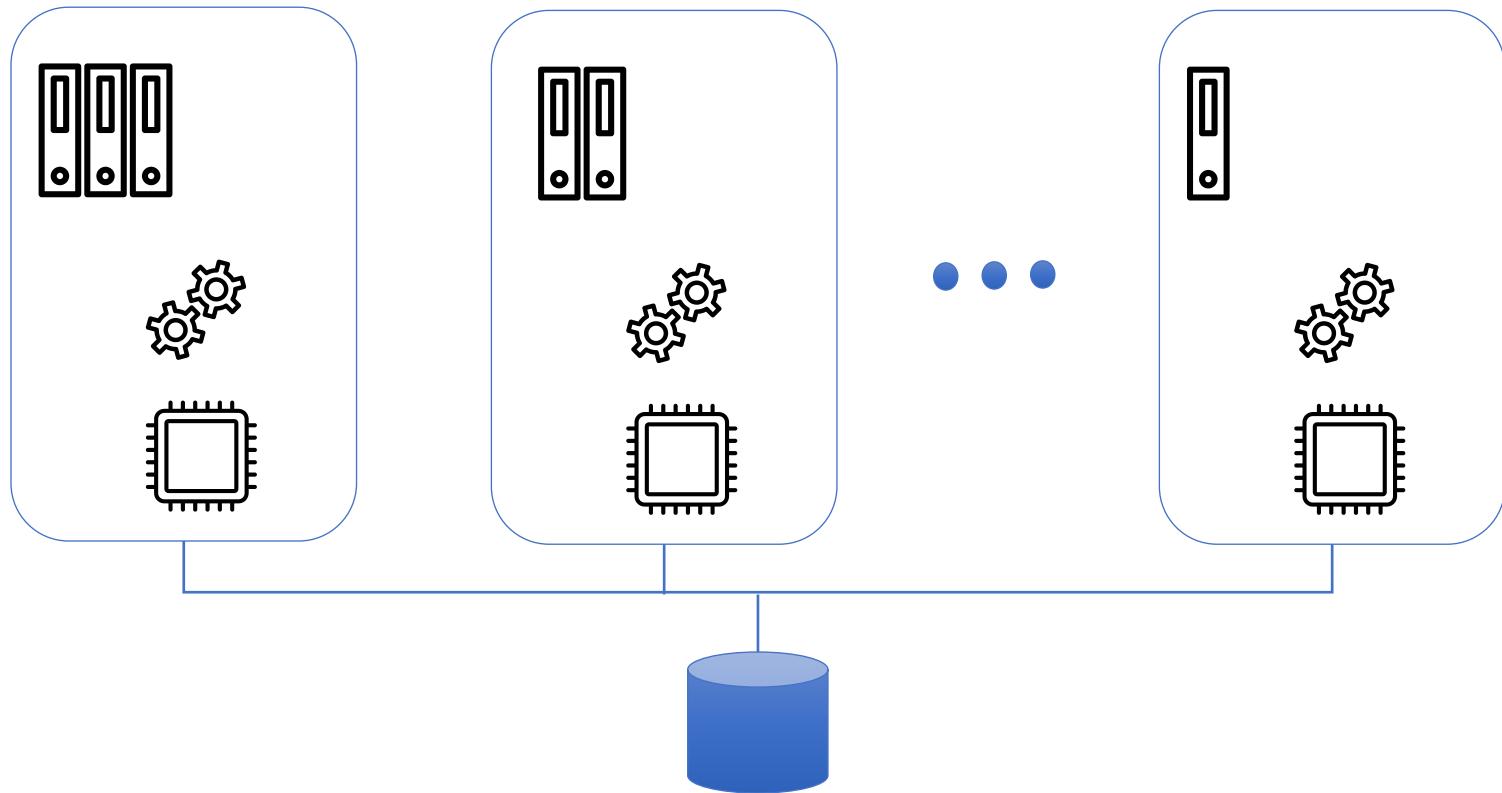
Load-balancing



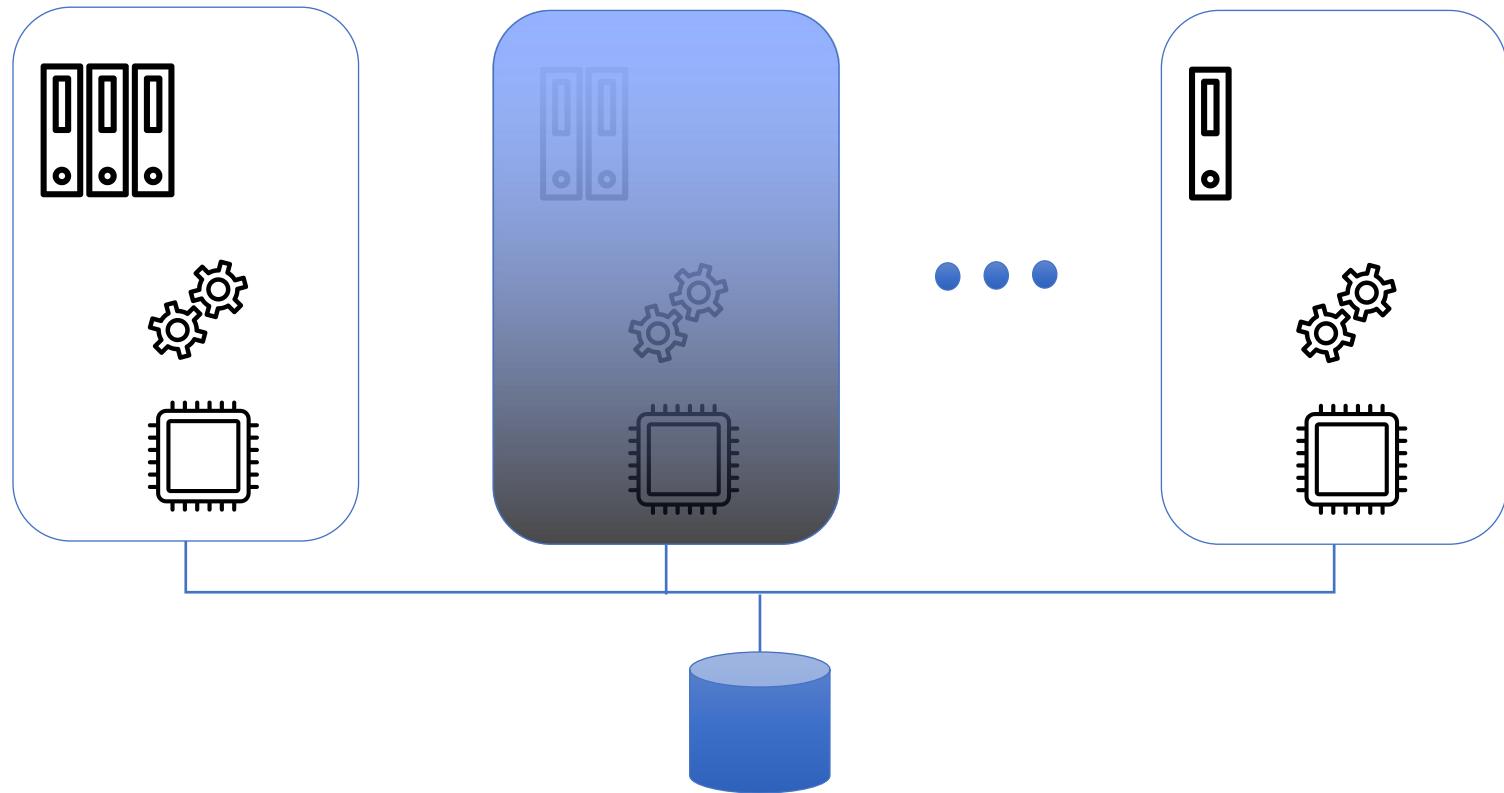
Load-balancing



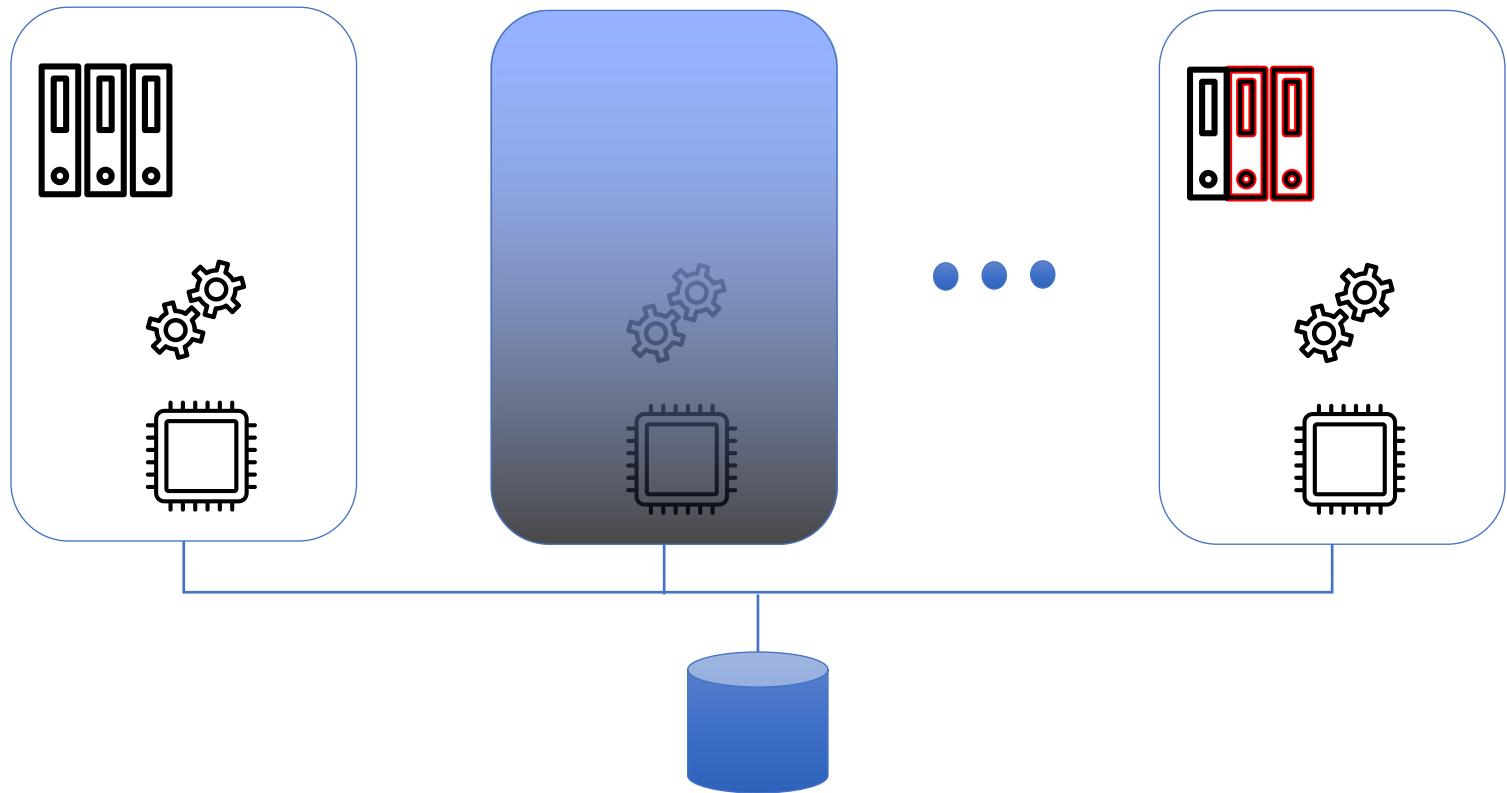
High-Availability (HA)



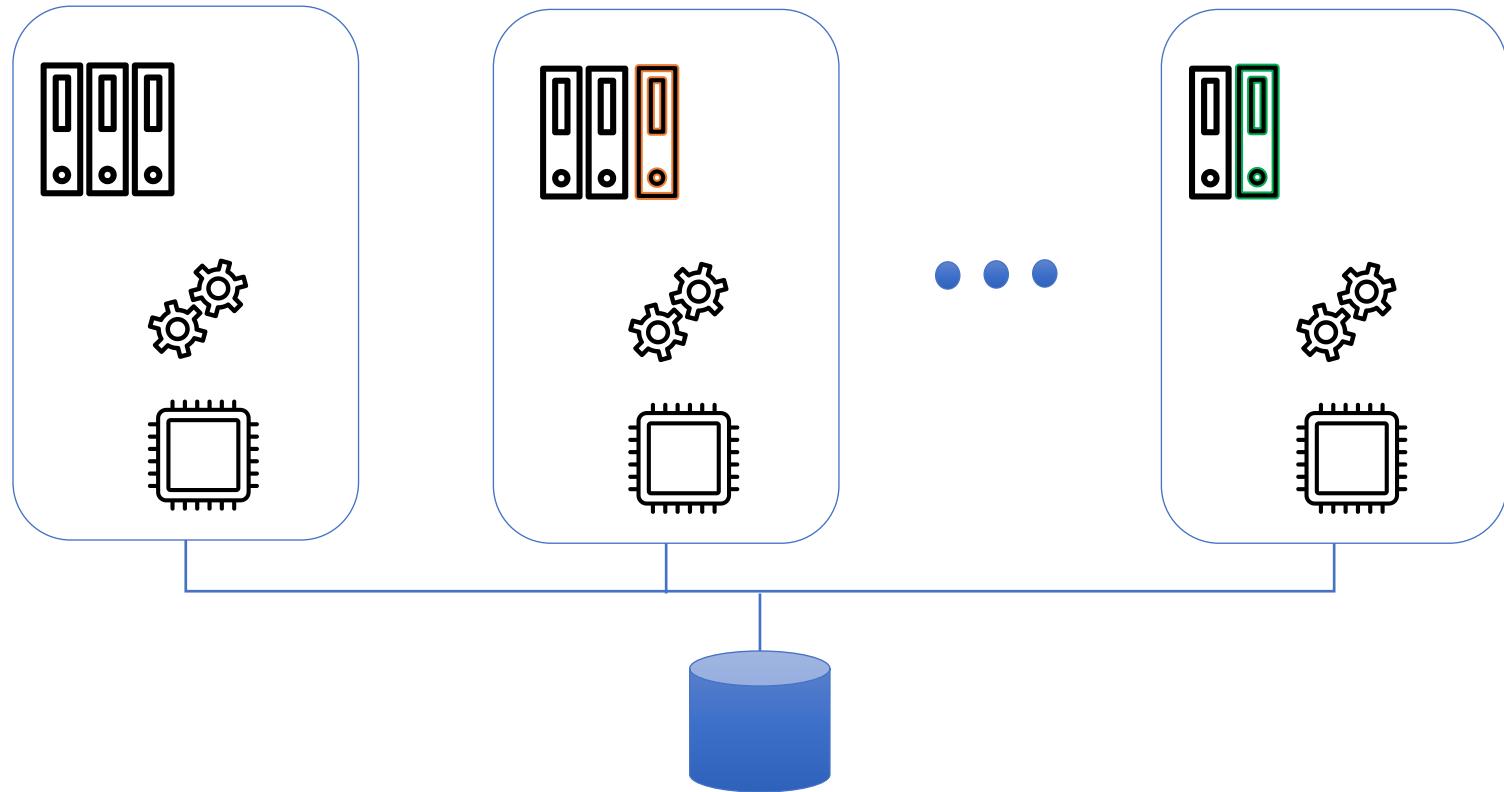
High-Availability (HA)



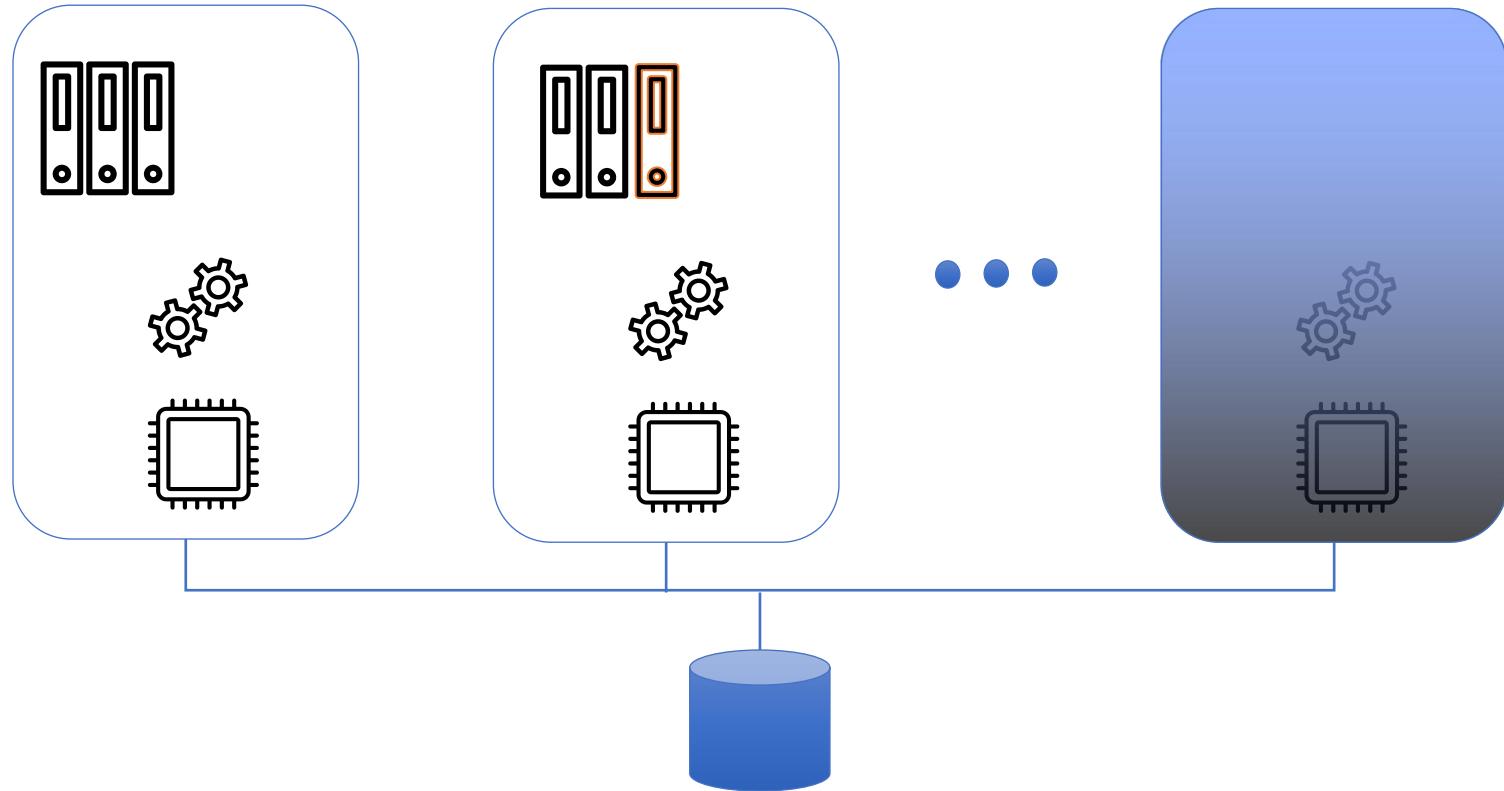
High-Availability (HA)



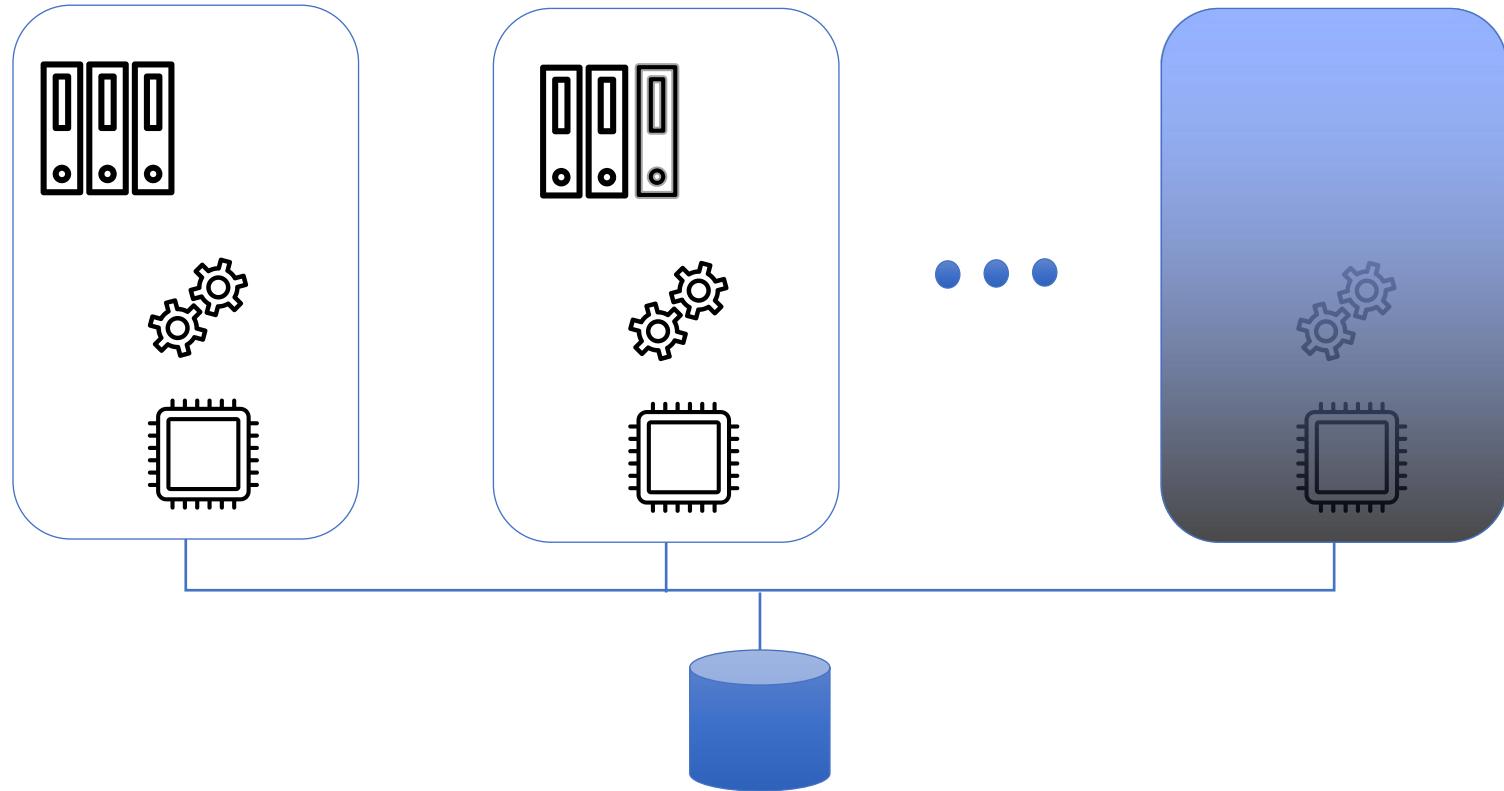
Fault-Tolerance (FT)



Fault-Tolerance (FT)

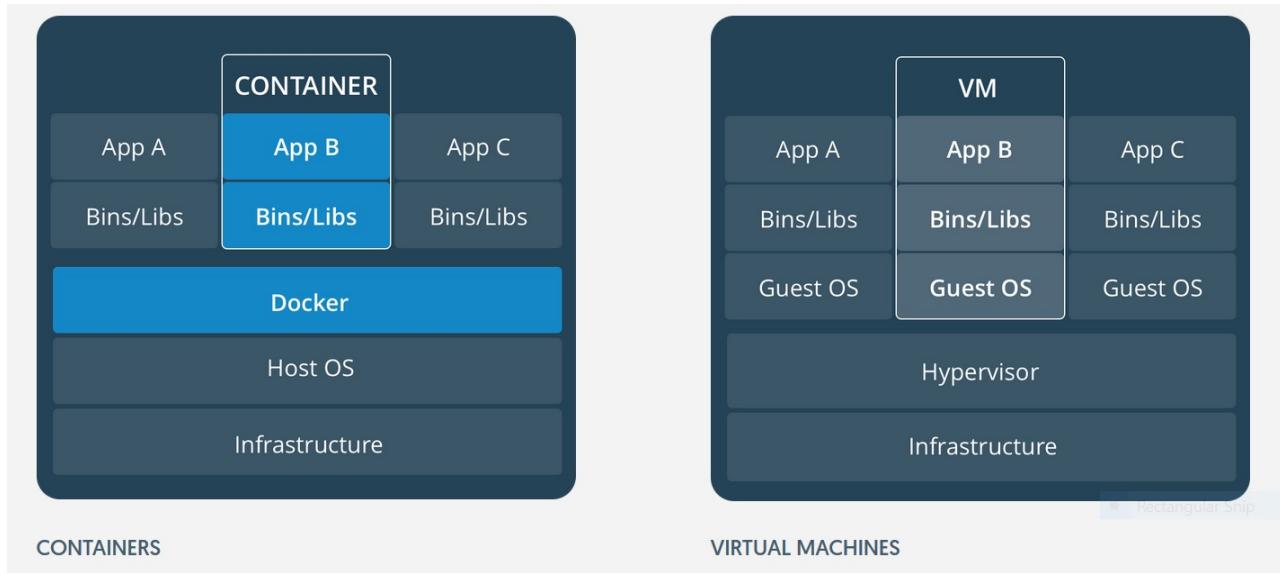


Fault-Tolerance (FT)





Containers





Full scale VMs vs Containers

- Hypervisors are usually more resource management efficient
- Virtualization clusters have HA & FT features
- Containers are faster to start than VMs
- Containers are usually micro-service orientated
- By default, containers work in a stateless mode
- Containers are portable
- Containers can also be used for local needs (dev environments, sandboxes,...)
- Many orchestration tools exist for containers (Kubernetes, Docker Swarm)
- VMs can be moved between physical host without being stopped
- Web services can be provided through docker images installed in VMs



Virtualization in a nutshell

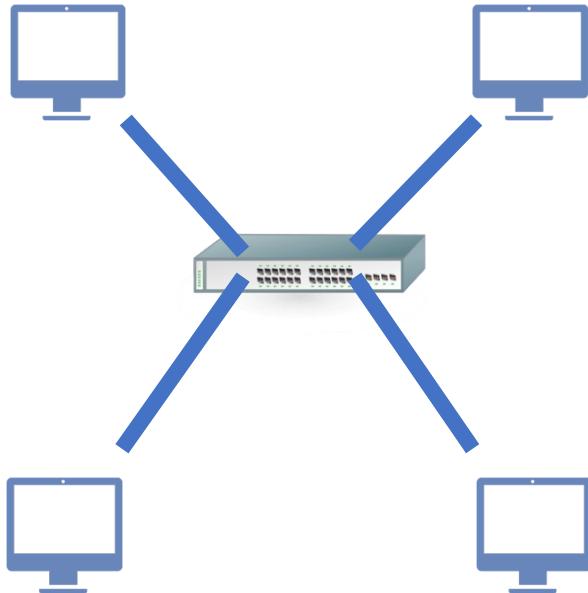
Network virtualization



Network Virtualization

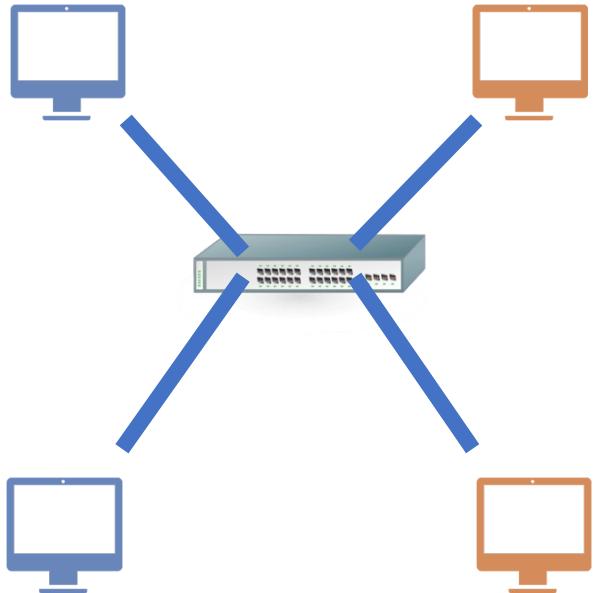
- Physically based on VLANs
- “Virtually” based on Software Defined Networks (SDN)

No VLANs: same L3 network



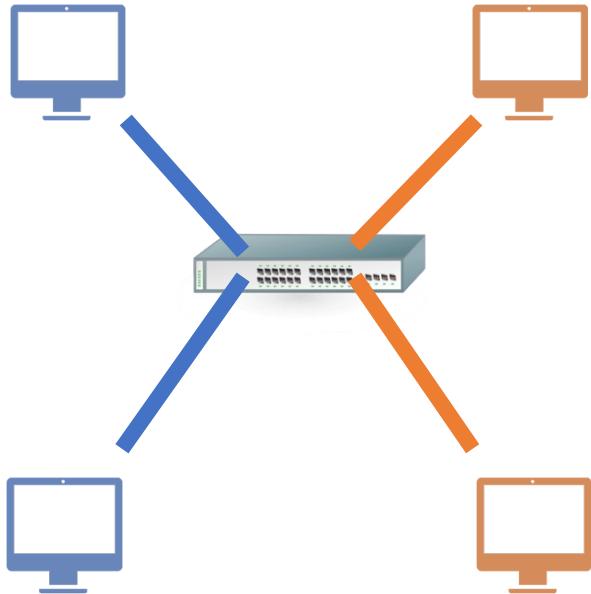
- All machines are
 - In the same L2 network
 - In the same L3 network
- Consequences:
 - They can all communicate
 - They are in the same (L2) broadcast domain

No VLANs: different L3 networks



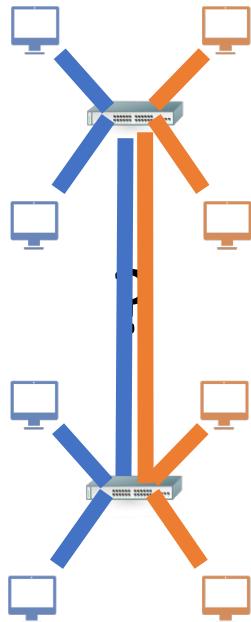
- Machines are
 - In the same L2 network
 - In different L3 networks
- Consequences:
 - They are separated at the L3 level
 - They are in the same (L2) broadcast domain
 - Isolation is only based on L3; no real security

Per-port VLANs



- Machines are
 - In different VLANs (i.e. L2 networks)
 - L3 configuration doesn't matter (unless communication between networks is required)
- Consequences:
 - There are different (L2) broadcast domain
 - Isolation is now based on L2

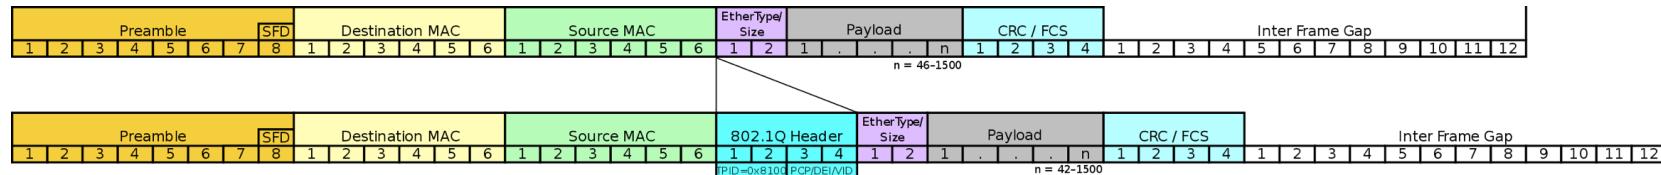
Per-port VLANs limitations



- How to connect different switches preserving VLAN configuration ?
- Answer:
 - Use dedicated ports on each switch
- Problem:
 - When the number of VLANs or the number of switches increases, more ports are "wasted" for the interconnection



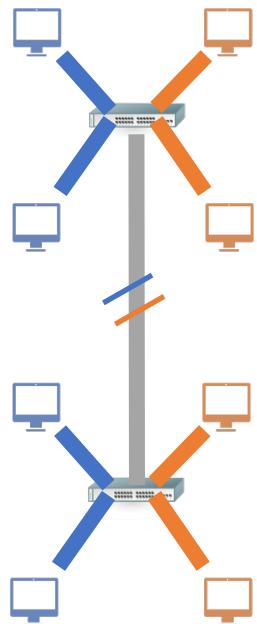
802.1Q



16 bits	3 bits	1 bit	12 bits
TPID 0x8100	TCI		
	Priority Code Point	Drop eligible Indicator	VLAN-ID 0-4095*

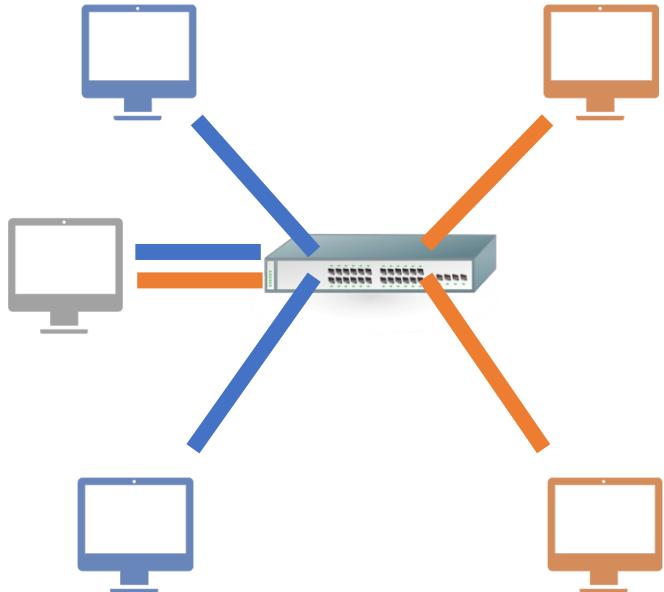
*: 0x000 and 0xFFFF are reserved

Solution through frame tagging



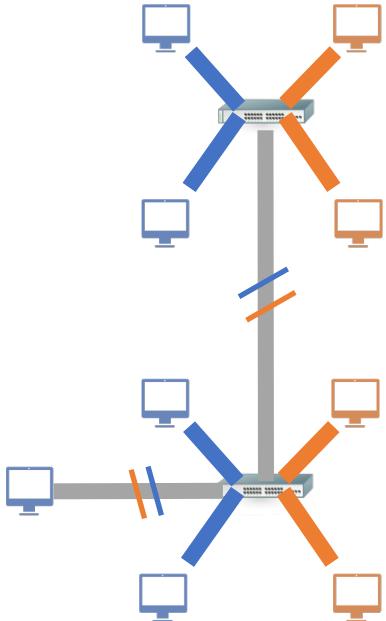
- Ethernet frames are modified to handle "tags" or VLAN-Ids
- Only one wire is required
- Switches have to understand tagging
- Bandwidth bottlenecks can appear

VLAN interconnection without tagging



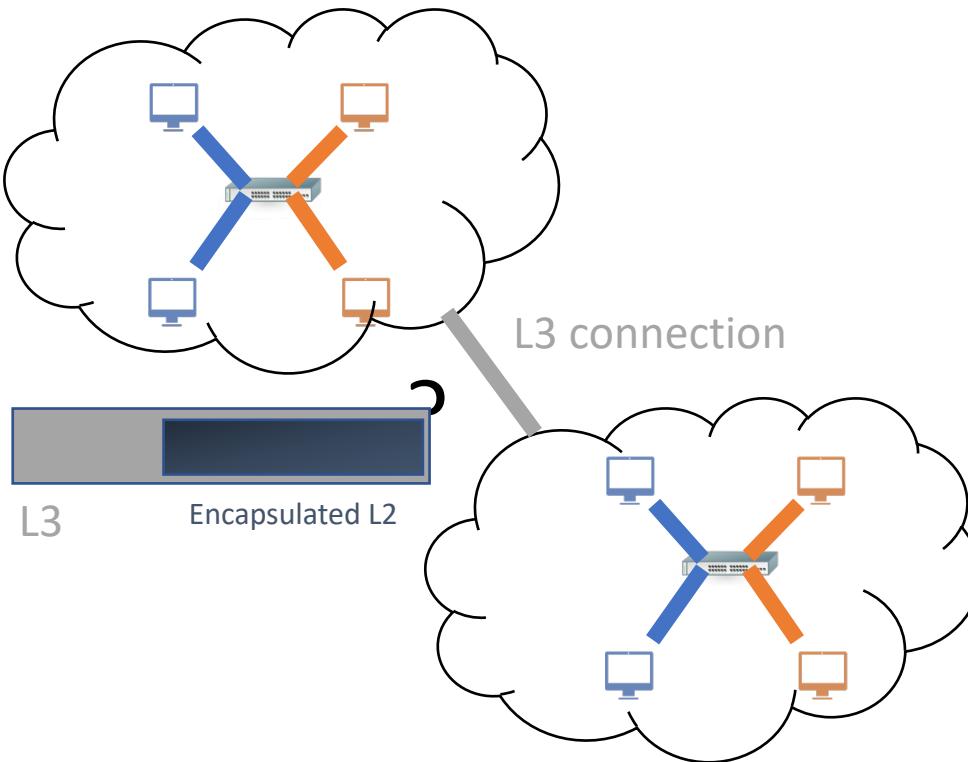
- How to interconnect two VLANs ?
- L2 levels are isolated
 - Interconnection at a higher level through routing
 - IP addressing plans must be different
- The same architecture can be used for services installed on both networks
- The router or server and the switch need as many physical cards or ports as there are VLANs

VLAN interconnection with tagging



- How to interconnect two VLANs ?
- L2 levels are isolated
 - Interconnection at a higher level through routing
 - IP addressing plans must be different
- The same architecture can be used for services installed on both networks
- The router or server has a single network card, and a single port is used on the switch
- The router or server needs to understand tagging

What about distant L2 networks ?



- Interconnection usually exists at the L3 level
- Answer:
 - Encapsulate the original L2 frame into an L3 packet
- Problem:
 - Interconnection points need to understand the encapsulation technique

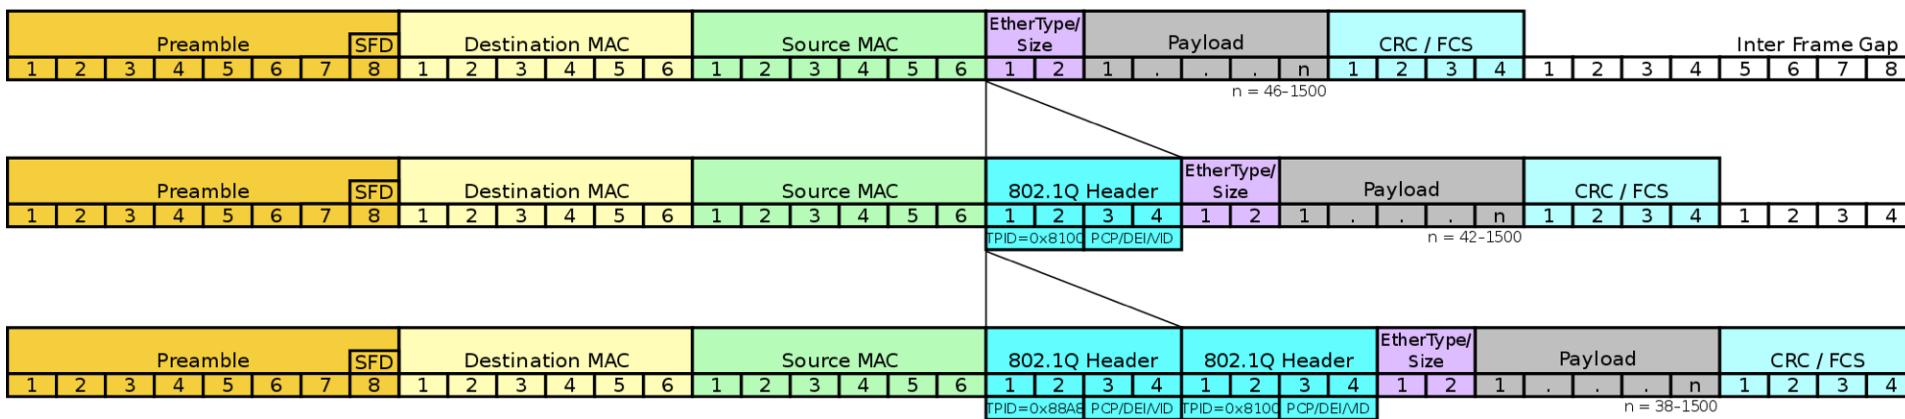


Other “lower level” approaches

- Problem:
 - Specific VLAN assignment: (theoretically) limited to 2^{12} VLANs
- Standard solutions:
 - Q-in-Q
 - VXLAN

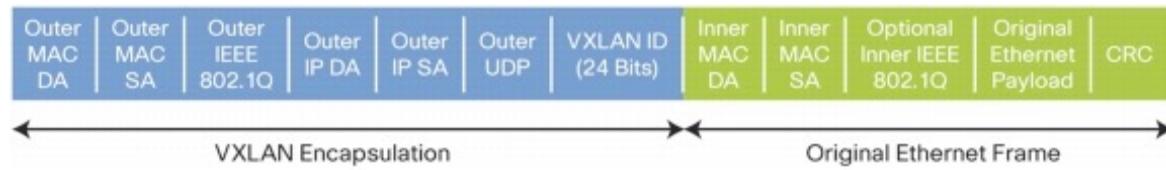


L2 Solution: 802.1ad (Q-in-Q)





L3 solution: VXLAN



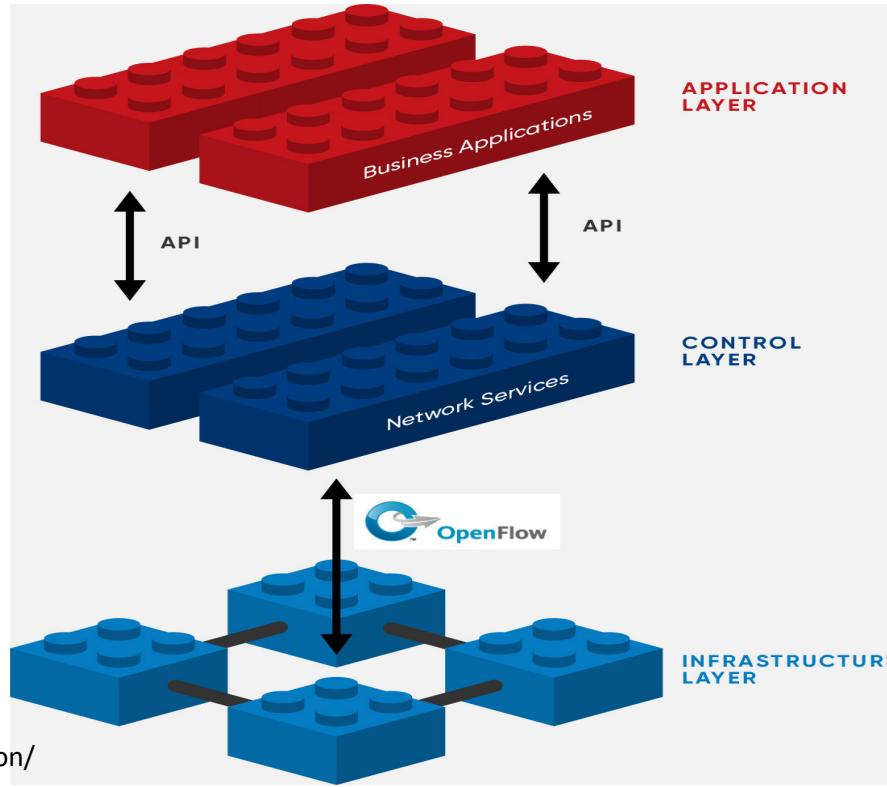


The SDN approach

- Goals:
 - Create programmable network architectures
 - Allows automation,
 - Speeds up deployment
 - Provides agility at the infrastructure level
 - Based on flow routing
 - Allows QoS rules



The SDN approach



Source: <https://opennetworking.org/sdn-definition/>



The SDN approach

- API enabled
- Provides Firewalling / Security groups
 - Applies rules down to the host from outside the host
- Security policies can be based on tags
- Some SDN solutions & implementations:
 - OpenFlow,
 - Openstack (Neutron),
 - VMWare NSX.



Virtualization in a nutshell

Storage virtualization



In computer science, **storage virtualization** is the process of presenting a logical view of the physical storage resources to a host computer system

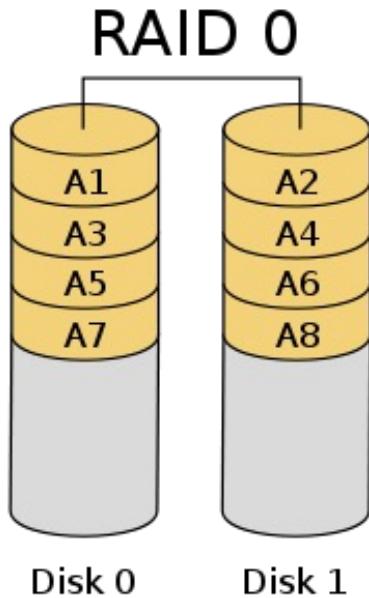


Understanding Storage Technologies

- At the physical level,
 - Example: RAID,
- Block Storage,
 - Examples: LVM, SAN
- File Storage,
 - Examples: NAS
- Object Storage.
 - Example: Ceph



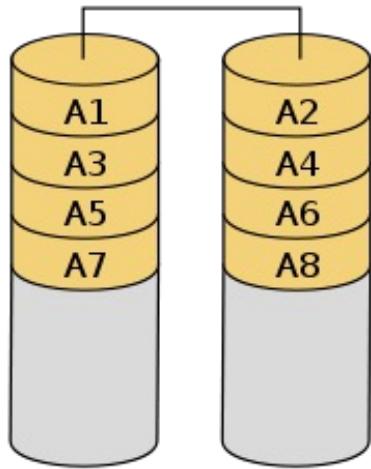
RAID





RAID

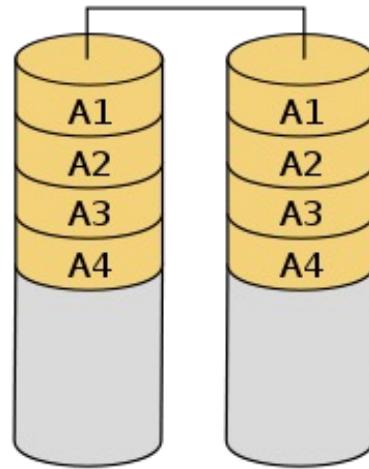
RAID 0



Disk 0

Disk 1

RAID 1



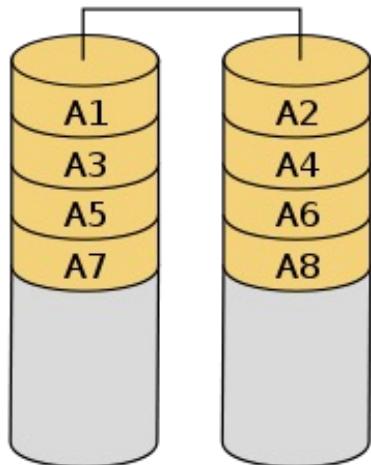
Disk 0

Disk 1

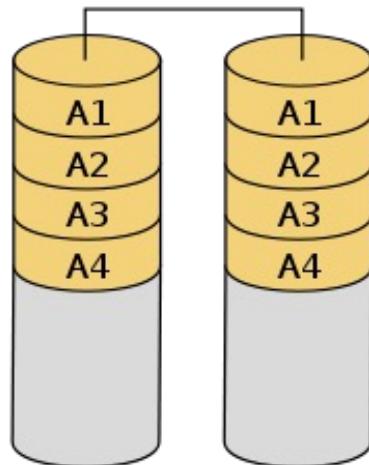


RAID

RAID 0

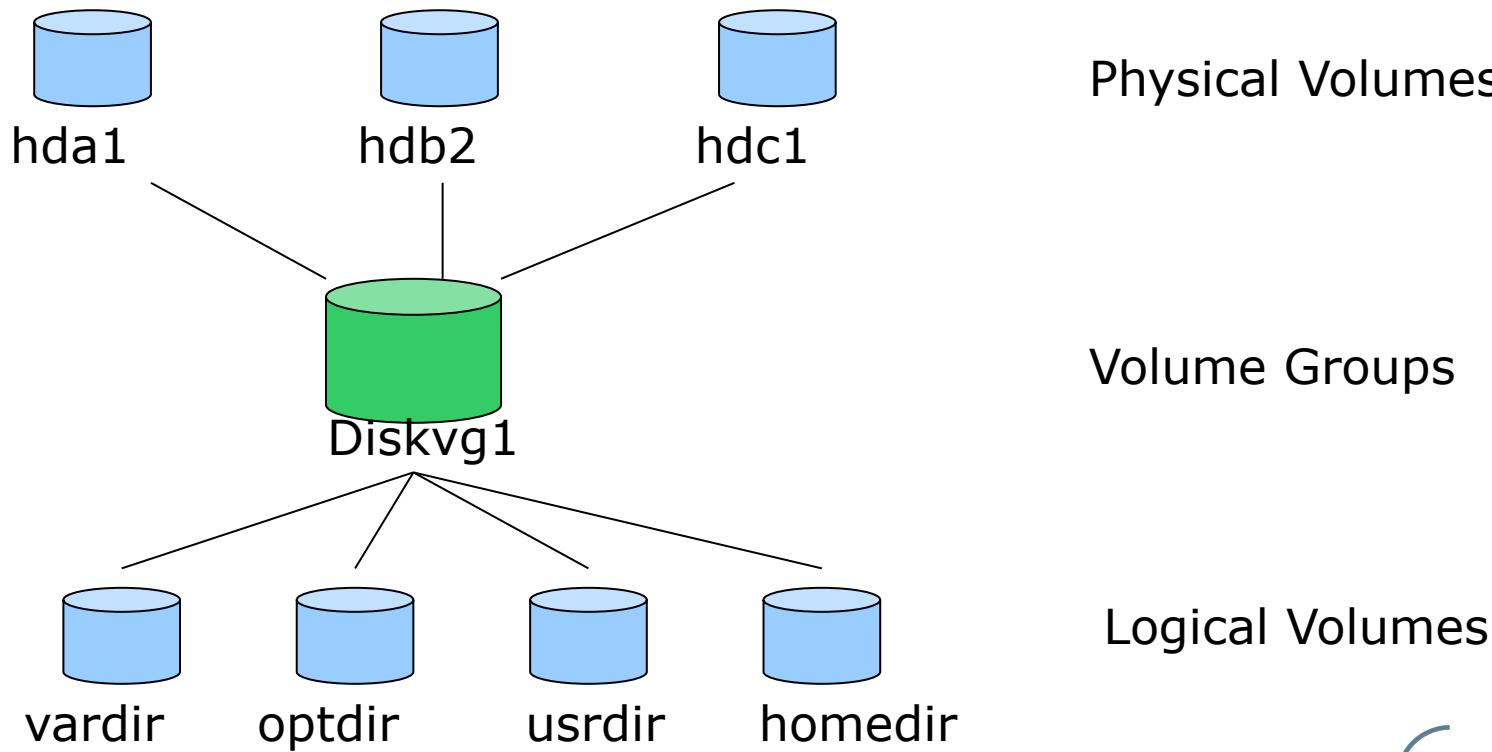


RAID 1



RAID 5







Network Attached Storage

- File sharing solutions based on IP
- Based on FILES
 - Client systems access files,
 - Files security is enforced at the server level,
 - The real file system is hidden to the client
- Standard solutions:
 - CIFS (or SMB)
 - NFS
 - AFP



Storage Area Network

- Based on BLOCKS
 - The client system “attaches” a remote disk,
 - The client system must install a filesystem,
 - The server controls the access to the shared device,
 - The server doesn’t use the filesystem installed by the client.
- Standard solutions:
 - Dedicated Fiber Channel network,
 - iSCSI on IP networks.



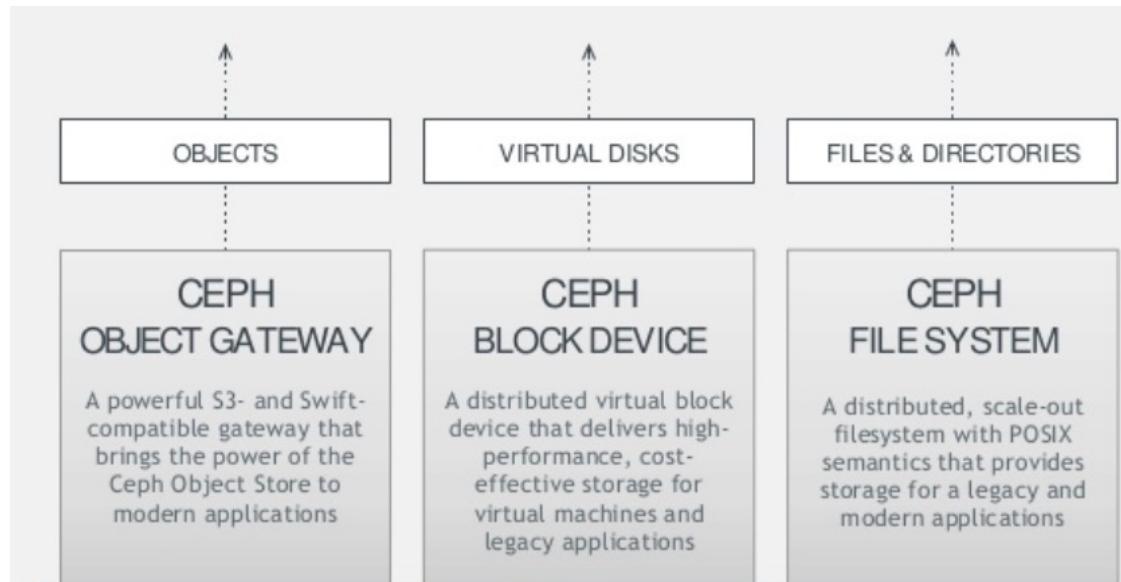
Object Storage

- Basics:
 - Objects are saved “as is” in a flat file system
 - Metadata are associated to objects identified by IDs
- Pros:
 - Scalability,
 - Customizable metadata,
 - High (sequential) throughput performance.
- Cons:
 - Full object manipulation,
 - Can’t easily be used as standard filesystems
- Standard solutions:
 - Ceph Object Storage, Swift, Amazon S3



A focus on Ceph

- Ceph is an opensource, unified, distributed storage system designed for performance, reliability and scalability.

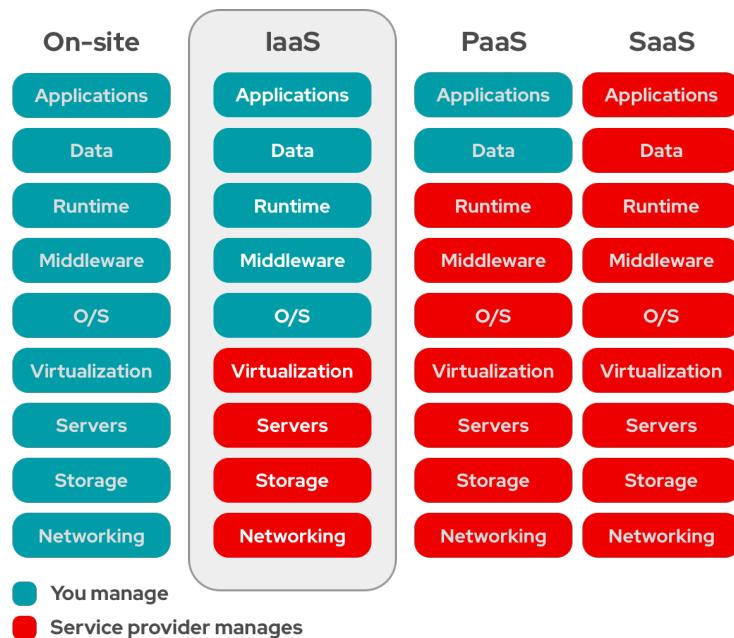




Back to Cloud Computing...



- Standard Cloud Services Models – SPI Stack
 - Software as a Service (SaaS)
 - Platform as a Service (PaaS)
 - Infrastructure as a Service (IaaS)

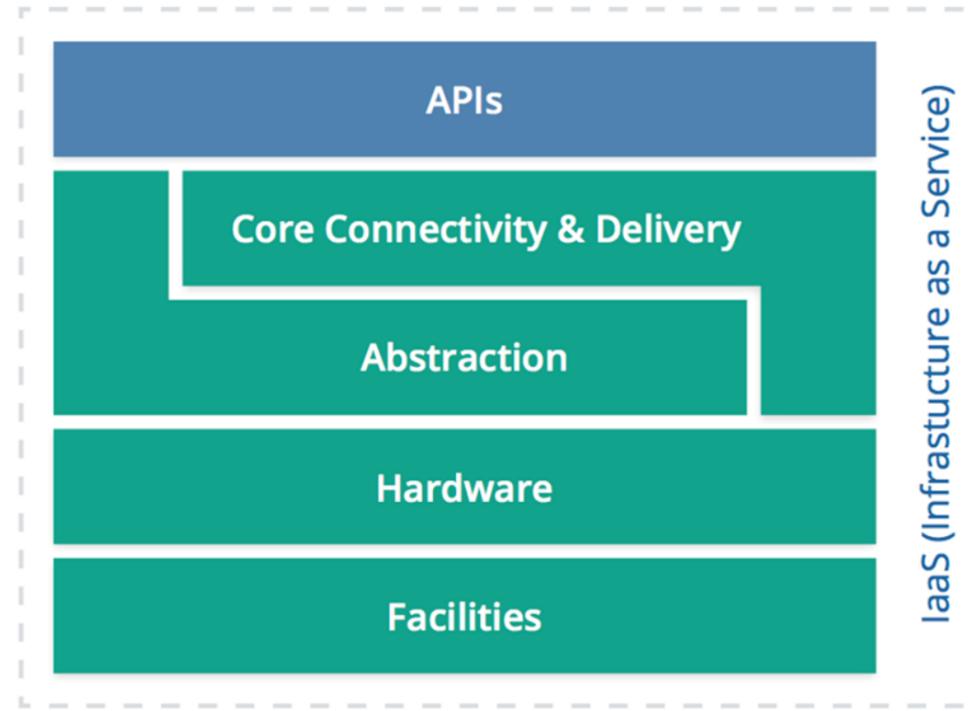


Source: <https://www.redhat.com/en/topics/cloud-computing/what-is-iaas>

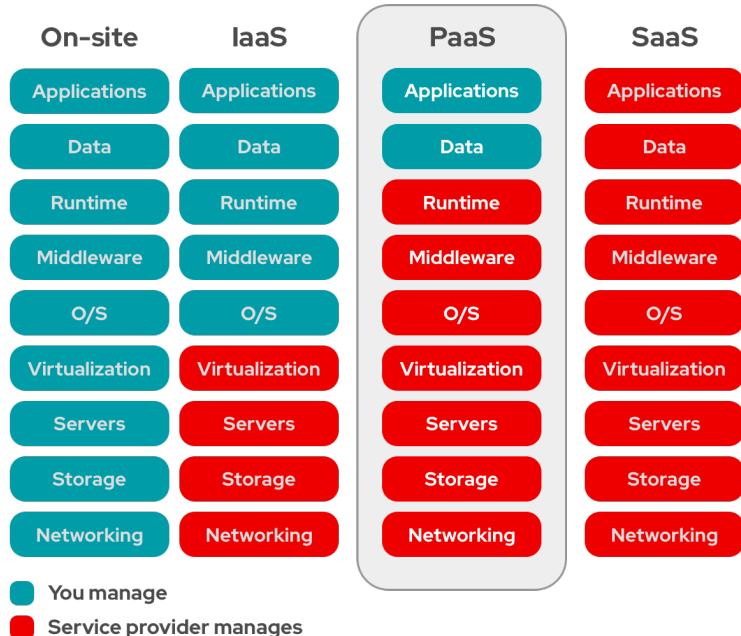


Examples:

- AWS EC2
- Google Compute Engine
- OVHCloud Public Cloud



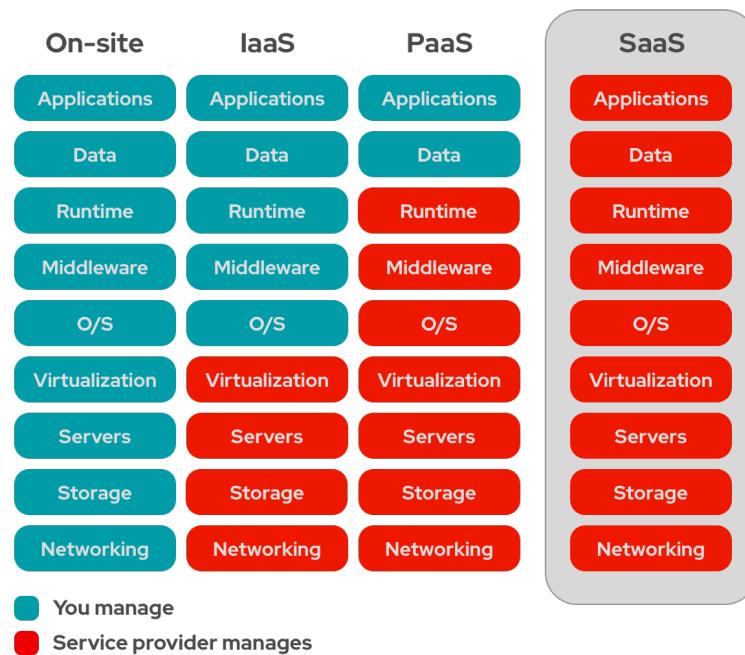
Source: CSA - Cloud Security Course 1.0 Netacad



Source: <https://www.redhat.com/en/topics/cloud-computing/what-is-paas>



- Examples:
 - AWS Beanstalk
 - AWS Lambda
 - Google App Engine
 - Microsoft Azure
 - OVHcloud Web PaaS
 - IBM Cloud Foundry



Source: <https://www.redhat.com/en/topics/cloud-computing/what-is-saas>



- Examples:
 - Google G Suite
 - Microsoft Office 365
 - Dropbox
 - Slack



Cloud Deployment Models

- The NIST model defines 4 types of Cloud models
 - Public Cloud,
 - Private Cloud,
 - Hybrid Cloud,
 - Community Cloud.



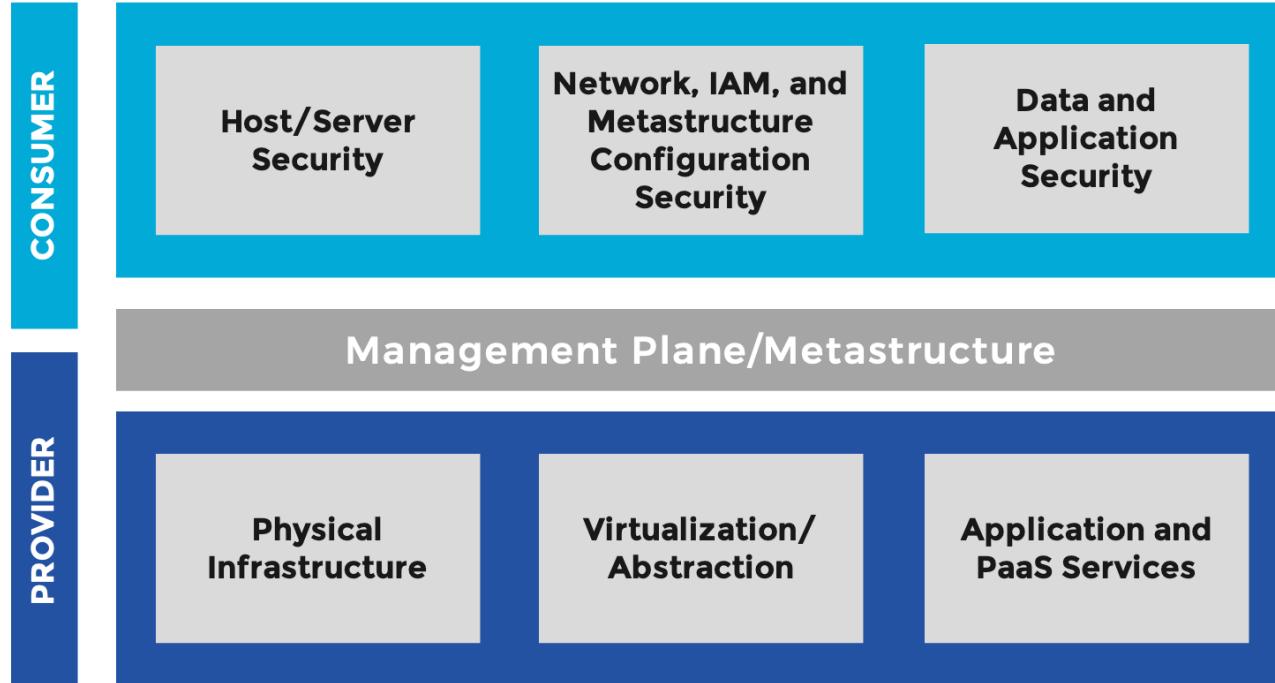
Cloud Deployment Models

	Infrastructure Managed By ¹	Infrastructure Owned By ²	Infrastructure Located ³	Accessible and Consumed By ⁴
Public	Third Party Provider	Third Party Provider	Off-Premise	Untrusted
Private/Community	Organization Third Party Provider	Organization Third Party Provider	On-Premise Off-Premise	Trusted
Hybrid	<u>Both</u> Organization & Third Party Provider	<u>Both</u> Organization & Third Party Provider	<u>Both</u> On-Premise & Off-Premise	Trusted & Untrusted

Source: CSA - Cloud Security Course 1.0 Netacad



Shared responsibility model



Source: CSA - Cloud Security Course 1.0 Netacad



Responsibility vs. Models



Source: CSA - Cloud Security Course 1.0 Netacad



- Case studies:
 - The OVHCloud offer
 - The AWS offer
- Global things to consider:
 - Pricing
 - Availability zones
 - Security features